

Hierarchical Posterior Sampling for LLM Multi-Agent Coordination with Provable Exponential Burn-in Improvement

Anonymous Authors

Affiliation
email

Abstract

We address centralised multi-agent coordination of n LLM agents with private personas $\theta_i \in \Theta_i$ and unknown dynamics. We propose two algorithms: **H-PSMG**, which samples per-agent persona marginals and a transition kernel from a factored Bayesian posterior at each episode and plays the CCE of the sampled world; and **H-PSMG⁺**, which augments H-PSMG with a type-discrimination bonus computable directly from the factored marginals. Both algorithms interface with arbitrary LLM backbones through three mechanistic primitives – persona-as-prompt, sampled-world CCE oracle, and persona-likelihood scoring – with no fine-tuning. Theoretical grounding (deferred to the appendix for full proofs): the joint posterior over (P^*, θ^*) exactly decouples across the kernel and per-agent persona marginals under three structural conditions, giving exponential-in- n storage savings and Bayesian regret $\tilde{O}(H^{3/2}\sqrt{SA_{\text{joint}}K} + H\sqrt{|\Theta_i|K})$, exponentially tighter than the prior-factorisation lower bound $\Omega(\sqrt{|\Theta|^n K})$ and near-tight against the matching minimax bound. The H-PSMG⁺ bonus cuts deep-trap burn-in regret from $\Theta(|\Theta_i|^n H)$ to $O(H)$, an exponential separation. We evaluate across three LLM substrates of increasing wildness on four backbones (DeepSeek-V3.2, GPT-5.4-nano, Kimi-K2.6, Llama-4-Maverick): (i) **HP-SPGG**, a controlled persona-coordination benchmark we construct, where the H-PSMG family beats the best published LLM-coordination baselines (ECON (Yang et al., 2025), A-ToM (Xu et al., 2026)) by 6–12 $\times$, a β -sweep reproduces all three regimes the theorem predicts (matched to per-backbone posterior concentration speed), and per-episode trajectories validate the $\tilde{O}(\sqrt{K})$ rate of H-PSMG against PSRL-NoType’s linear-in- K growth with 11–30 $\times$ separation at $K=20$; (ii) **DeepMind Concordia** (Vezhnevets et al., 2023), where H-PSMG⁺ wins or ties oracle on 16/16 configurations; (iii) **SOTOPIA-**

Hard (Zhou et al., 2024), where the prior-factorisation assumption is violated and the substrate is structurally out of scope; the substrate-adapted H-PSMG⁺ transfers partially, leading on the largest scenario family (*craigslist_bargains*, 5/16 wins) and remaining within 0.1 of the best alternative on every backbone, with the exponential dominance of HP-SPGG and Concordia gracefully vanishing as the theory predicts.

1 Introduction

This paper is motivated by an empirical observation across four LLM backbones (DeepSeek-V3.2, GPT-5.4-nano, Kimi-K2.6, Llama-4-Maverick) and three multi-agent coordination benchmarks: in heterogeneous-type settings where each agent has a private persona drawn from a small archetype set, type-aware Bayesian methods attain $\tilde{O}(\sqrt{K})$ cumulative regret while strong type-agnostic LLM-coordination baselines (PSRL-NoType, prompted-greedy) accumulate $\Omega(K)$ linear regret on the same instances – an 11-30 $\times$ cumulative separation at horizon $K=20$. We provide the algorithm and theory that explain and exploit this gap.

Multi-agent reinforcement learning in environments where agents differ in preferences or roles arises in cooperative AI, mechanism design, and economic simulation (Zhang et al., 2021; Leibo et al., 2017; Perolat et al., 2017). The Markov-game model (Shapley, 1953; Littman, 1994) captures sequential interaction but treats agents as exchangeable in type. In practice agents have private latent attributes – foremost, their reward functions – and reasoning under this uncertainty is the defining feature of Bayesian games (Harsanyi, 1967, 1968), yet remains under-developed in the sequential MG setting with explicit regret guarantees.

Posterior sampling for RL (PSRL) (Strens, 2000; Osband et al., 2013; Agrawal and Jia, 2017; Osband and Van Roy, 2017) has produced near-optimal regret in single-agent MDPs. Extensions to two-player *zero-sum* MGs were given by Jafarnia-Jahromi et al. (2021); Xiong et al. (2022); Yueh et al. (2026), and cooperative parallel-MDP exploration by Hsu et al. (2024). None addresses the natural cooperative-AI regime where agents are het-

erogeneous in reward type *and* the learner must infer both transitions and the joint type profile.

Contributions. (1) **Algorithm family.** We propose two algorithms for centralised multi-agent LLM coordination under unknown personas and dynamics: *H-PSMG* (Algorithm 1), which samples per-agent persona marginals and a transition kernel from a factored posterior at each episode and plays the CCE of the sampled world; and *H-PSMG⁺* (Algorithm 2), which augments H-PSMG with a *type-discrimination bonus* computable only from the factored marginals. Both algorithms interface with arbitrary LLM backbones via three mechanistic primitives – persona-as-prompt, sampled-world CCE oracle, and persona-likelihood scoring – with no LLM fine-tuning. (2) **Theoretical grounding.** Under three structural conditions (transition independence, own-type reward locality, factored prior), the joint posterior over (P^*, θ^*) exactly decouples across the transition kernel and per-agent persona marginals at every episode (Posterior Decoupling Theorem, Appendix B). This yields storage savings exponential in n and a Bayesian regret bound $\tilde{O}(H^{3/2}\sqrt{SA_{\text{joint}}K} + H\sqrt{|\Theta_i|K})$ for H-PSMG, exponentially tighter than the prior-factorisation lower bound $\Omega(\sqrt{|\Theta|^n K})$ in the absence of factorisation, and near-tight against the minimax lower bound. For H-PSMG⁺, the bonus reduces burn-in regret on a deep-trap family from $\Theta(|\Theta_i|^n H)$ to $O(H)$, an $\Omega(|\Theta_i|^n)$ exponential separation (Burn-in Separation Theorem E.4). (3) **LLM-substrate evaluation across three benchmarks of increasing wildness.** We construct *HP-SPGG*, a controlled $n=3$ persona-coordination benchmark satisfying all six HT-MG-BC assumptions; report cross-model regret on four backbones (DeepSeek-V3.2, GPT-5.4-nano, Kimi-K2.6, Llama-4-Maverick) against a 16-method suite that includes the best published LLM-coordination baselines (ECON (Yang et al., 2025), A-ToM (Xu et al., 2026)); a β -sweep that reproduces all three regimes the theorem predicts (matched to per-backbone posterior concentration speed); and four theory-driven experiments (E1–E5) confirming the predicted scaling in $|\Theta_i|$, n , and the rate separation against type-agnostic baselines (11–30 $\times$ at $K=20$). On the independent DeepMind Concoridia benchmark (Vezhnevets et al., 2023) H-PSMG⁺ wins or ties oracle on 16/16 configurations. On SOTOPIA-Hard (Zhou et al., 2024), where the prior-factorisation assumption is violated and the substrate is structurally out of scope ($n=2$ dyadic, open-ended natural-language actions, non-factored persona prior), the substrate-adapted variant transfers partially – leading on the largest scenario family (*craigslis bargains*, 5/16 wins) and remaining within 0.1 of the best alternative on every backbone – demonstrating graceful degradation rather than catastrophic failure when the theory’s preconditions

are out of scope.

2 Related Work

Posterior sampling. PSRL (Strens, 2000; Osband et al., 2013; Agrawal and Jia, 2017; Ouyang et al., 2017) attains near-optimal Bayesian regret in single-agent MDPs, with the information-theoretic framework of Russo and Van Roy (2014, 2016) underpinning our type-inference analysis. In Markov games, posterior sampling has so far been restricted to zero-sum: Jafarnia-Jahromi et al. (2021) for average-reward, Xiong et al. (2022) with function approximation, and Yueh et al. (2026) for episodic finite-horizon. All assume known type structure.

Latent-type and heterogeneous MARL. Kwon et al. (2021) prove that learning latent MDPs is exponentially hard without identifiability; Tian et al. (2024) extend hardness to adaptive adversaries. Albrecht and Ramamoorthy (2013); Albrecht et al. (2019) introduce Harsanyi-Bellman Ad-hoc Coordination, a Bayesian opponent-type method with asymptotic convergence but no finite-time regret. Our decoupling theorem and $\tilde{O}(\sqrt{K})$ bound provide the finite-sample complement.

General-sum equilibrium learning and PSRO. A line on V-learning and its variants establishes sample complexity for CCE in general-sum MGs (Jin et al., 2021; Song et al., 2022; Mao and Başar, 2023; Erez et al., 2023; Cui et al., 2025); the PSRO family (Lanctot et al., 2017; Muller et al., 2020; McAleer et al., 2020; Marris et al., 2021; McAleer et al., 2021, 2022; Dinh et al., 2022; Zhou et al., 2022; Yao et al., 2023) targets equilibrium of an empirical game via population learning. Neither line maintains a Bayesian posterior over latent payoff structure; our framework is complementary and PSRO-style oracles can be plugged into the inner CCE step of H-PSMG without affecting the decoupling result.

Structured Bayesian RL. Mutti et al. (2024) factor dynamics via causal priors; Arumugam and Griffiths (2026) revive PSRL for LLM agents. Neither factors the prior across agent types.

3 H-PSMG and H-PSMG⁺: Hierarchical Posterior Sampling for LLM Coordination

3.1 Setting

A centralised coordinator oversees n LLM agents over K episodes of horizon H . Each agent i has a private persona $\theta_i \in \Theta_i$ drawn from a discrete archetype set; the joint persona profile $\theta^* = (\theta_1^*, \dots, \theta_n^*)$ is sampled once and held fixed across episodes. At each step the coordinator observes the joint state, samples a joint action via a posterior over (P^*, θ^*) , executes it through the LLM agents, and records rewards $r_i(s, a, \theta_i^*) \in [0, 1]$. The perfor-

Algorithm 1 H-PSMG

Input: priors $P_0^P, \{P_0^{\Theta,i}\}$; horizon H ; episodes K ; CCE oracle.

- 1: Initialise $\mu_1^P \leftarrow P_0^P, \mu_1^{\Theta,i} \leftarrow P_0^{\Theta,i}$ for all i .
 - 2: **for** $k = 1, \dots, K$ **do**
 - 3: Sample $\hat{P}_k \sim \mu_k^P$ and $\hat{\theta}_k^i \sim \mu_k^{\Theta,i}$ independently.
 - 4: $\pi_k \leftarrow \text{CCE}(\hat{P}_k, \hat{\theta}_k)$; execute π_k , observe τ_k .
 - 5: Bayes-update each marginal independently from τ_k .
 - 6: **end for**
-

mance criterion is the Bayesian regret $\text{Reg}_i^B(K) = \mathbb{E}_{(P^*, \theta^*) \sim P_0} \sum_k (V_i^{\pi^{\text{CCE}}(P^*, \theta^*)}(s_1^k; \theta_i^*) - V_i^{\pi_k}(s_1^k; \theta_i^*))$, where π^{CCE} is the coarse correlated equilibrium of the true world. The formal HT-MG-BC tuple, the six structural assumptions (TI / RL / PF / TID / ISI / CSR), and the necessity counter-examples are deferred to Appendix A. The remainder of this section presents the two algorithms in the H-PSMG family together with the LLM-substrate implementation.

3.2 H-PSMG: factored posterior sampling

Why factor? The joint posterior over θ has $\prod_i |\Theta_i|$ entries, exponential in n . We instead maintain a separate marginal $\mu_k^{\Theta,i}$ for each agent, plus a single transition posterior μ_k^P . This reduces storage to $O(n|\Theta_i|)$; the reduction is exact rather than approximate, by virtue of the following structural result.

Theory idea. Under transition independence of the persona profile (TI), own-type reward locality (RL), and a factored prior (PF), the joint posterior over (P^*, θ^*) exactly equals $\mu_k^P(P) \prod_i \mu_k^{\Theta,i}(\theta_i)$ at every episode (*Posterior Decoupling Theorem*; statement and induction proof in Appendix B). The decoupling also implies pathwise equivalence with an idealised Joint-PSRL under coupled randomness, so all guarantees and lower bounds transfer between the two while H-PSMG enjoys the $O(|\Theta_i|^n / (n|\Theta_i|))$ storage advantage.

LLM-based implementation. The factored posterior interfaces with an LLM substrate through three mechanistic primitives, each invoked once per episode. (1) *Persona as prompt template.* Each $\theta_i \in \Theta_i$ is realised as a system-prompt block (e.g. *altruistic builder*, *conditional cooperator*, *strategic hedger*, *free rider* in HP-SPGG); the LLM agent for slot i is instantiated with the prompt corresponding to its true θ_i^* , hidden from the coordinator. (2) *Trajectory generation under the sampled world.* The CCE oracle is computed by mechanistic backward induction on the tabular DP induced by the sampled world $(\hat{P}_k, \hat{\theta}_k)$; the resulting joint policy π_k

is then executed by querying each LLM agent for an action at each step. (3) *Persona likelihood for Bayes update.* Given the observed action sequence $(a_h^k)_h$ of agent i , the persona likelihood $L_{\Theta,i}(\tau_k; \theta_i)$ is computed by re-evaluating each emitted action under the persona-prompted LLM and reading off the next-token probability of the action token under each candidate $\theta_i \in \Theta_i$; per-step probabilities are multiplied to obtain $L_{\Theta,i}$, and the marginal is updated by Bayes' rule. The transition posterior μ^P is updated by standard Dirichlet conjugacy on the observed (s, a, s') transitions. All three primitives are LLM-substrate-agnostic and back-end interchangeable across the four backbones (DeepSeek-V3.2, GPT-5.4-nano, Kimi-K2.6, Llama-4-Maverick) in our experiments.

3.3 H-PSMG⁺: type-discrimination bonus

Failure mode of vanilla PSRL on HT-MGs.

When the sampled persona profile $\hat{\theta}_k$ produces a state in which all candidate types predict identical rewards, observations carry no information about types and the marginal stays uniform. The CCE oracle resolves the tied Q^* arbitrarily; an unlucky pattern of ties keeps the rollout in low-information states for $\Omega(|\Theta_i|^n)$ episodes before a chance success draw (the *deep-trap family*, formally analysed in Appendix J). H-PSMG⁺ eliminates this failure mode with a *type-discrimination bonus* computed directly from the factored marginals,

$$D_k(s, a) := \sum_{i=1}^n \mathbb{E}_{\theta, \theta' \sim \mu_k^{\Theta,i}} [r_i(s, a, \theta) - r_i(s, a, \theta')], \quad (1)$$

which is added with weight $\beta > 0$ to the Q -function in backward induction: $Q_h^+(s, a) \leftarrow Q_h(s, a) + \beta D_k(s, a)$. The bonus is non-zero precisely at state-action pairs where two posterior-sampled types would disagree on the reward, biasing the policy toward type-discriminative states.

Theory idea. On the deep-trap family, the bonus reduces expected burn-in regret from $\Theta(|\Theta_i|^n H)$ to $O(H)$ for any $\beta > 0$ – an $\Omega(|\Theta_i|^n)$ separation, exponential in the number of agents (*Burn-in Separation Theorem*, stated formally in §J; backward-induction proof in Appendix J).

LLM-based implementation. The bonus is computed entirely from the factored marginals at $O(n|\Theta_i|^2)$ cost per state-action pair, with *no additional LLM queries* beyond those required for H-PSMG. After sampling the world, the coordinator (i) evaluates Eq. (1) on the small tabular state-action grid, (ii) runs mechanistic backward induction with the augmented Q -function $Q_h^+(s, a) = Q_h(s, a) + \beta D_k(s, a)$, and (iii) executes the resulting policy via the same LLM-querying loop as H-PSMG. The bonus computation is therefore mechanistic and substrate-agnostic; only the

Algorithm 2 H-PSMG⁺

Input: as in Alg. 1, plus $\beta > 0$.

- 1: Initialise as in Alg. 1.
 - 2: **for** $k = 1, \dots, K$ **do**
 - 3: Sample $\hat{P}_k, \hat{\theta}_i^k$ as in Alg. 1; compute D_k via Eq. (1).
 - 4: Solve CCE with $Q_h^+(s, a) += \beta D_k(s, a)$ in backward induction.
 - 5: Execute, observe τ_k , Bayes-update marginals.
 - 6: **end for**
-

policy execution and the persona-likelihood scoring involve LLM calls. The single hyperparameter β is fixed at $\beta=0.25$ across all four backbones and substrates in our experiments and is not re-tuned per backbone.

The bonus is (i) *decoupling-driven* – it requires per-agent marginals, which Joint-PSRL recovers only by marginalising the joint posterior at $O(|\Theta|^n)$ cost; (ii) *self-decaying* – as $\mu_k^{\Theta, i}$ concentrates, $D_k \rightarrow 0$, so the bonus does not contaminate the asymptotic rate; (iii) *tie-resolving* – on trap families, $D_k > 0$ at type-discriminative states and propagates backward to break ties. Under matched RNG with $\beta=0$, H-PSMG⁺ is bit-identical to H-PSMG, which by the Posterior Decoupling Theorem is in turn bit-identical to Joint-PSRL.

Position vs prior PSRL-with-LLM work. H-PSMG generalises classical PSRL (Strens, 2000; Osband et al., 2013) from single-agent MDPs to heterogeneous-type multi-agent Markov games. The key technical novelty over the concurrent LLM-PSRL implementation of Arumugam and Grifiths (2025) – which runs vanilla PSRL with three LLM subroutines (posterior, sampling, policy) – is the Posterior Decoupling Theorem, which converts an $O(|\Theta|^n)$ -storage Joint-PSRL scheme into an $O(n|\Theta_i|)$ factored scheme *without statistical penalty*. H-PSMG⁺ further bypasses the burn-in regret bottleneck of vanilla PSRL on deep-trap families through the type-discrimination bonus.

What we can prove (informal). The H-PSMG family enjoys three regret guarantees, stated formally with proofs in Appendix B–E.5. (P1) *Storage and rate.* H-PSMG achieves Bayesian regret $\tilde{O}(H^{3/2}\sqrt{SA_{\text{joint}}K} + H\sqrt{|\Theta_i|K})$ at $O(n|\Theta_i|)$ storage; without prior factorisation, no algorithm beats $\Omega(\sqrt{|\Theta|^nK})$ (exponential gap in n), and the matching minimax bound is $\Omega(\sqrt{H^3SA_{\text{joint}}K} + \sqrt{|\Theta_i|HK})$ (near-tight modulo $\sqrt{H}$). (P2) *Rate separation.* On an explicit exploration-trap instance, H-PSMG attains $\tilde{O}(\sqrt{K})$ while MAP-Type-Greedy and PSRL-NoType both incur $\Omega(K)$ linear regret. (P3) *Exponential burn-in improvement.* On the deep-trap family $\mathcal{G}_m^{\text{trap}}$ parameterised by per-

agent type-space size m and agent count n , the expected burn-in regret of H-PSMG and Joint-PSRL is $\Theta(m^n H)$ while H-PSMG⁺ achieves $O(H)$ for any $\beta > 0$; the separation is $\Omega(m^n)$, exponential in n – which is the central justification for the bonus. Empirical signatures of all three are reported in §4.2 (Figures 3, 2, 4).

4 LLM Experiments

We evaluate on three benchmarks of increasing wildness: HP-SPGG, a controlled LLM-MAS substrate we construct that satisfies all six HT-MG-BC assumptions by construction; the public Concordia substrates from the NeurIPS 2024 Cooperative AI Contest; and SOTOPIA-Hard, whose persona structure violates PF. All experiments use four LLM backbones spanning vendors, capabilities, and weight access: *DeepSeek-V3.2* (open-weight MoE), *GPT-5.4-nano-20260317* (small closed), *Kimi-K2.6* (mid open), and *Llama-4-Maverick-17B-128E FP8* (large open).

Empirical validation framework. We frame claims about regret *rate* as testable predictions of (P1)–(P3), following the protocol of Park et al. (2025) adapted for the multi-agent setting. (i) *Sublinearity test for (P2):* we fit $\log \text{Reg}(t) = \hat{\alpha} \log t + \beta$ on per-episode cumulative regret and report $\hat{\alpha}$ with a one-sided test against the null $\hat{\alpha} \geq 1$ (linear regret); $\hat{\alpha} < 1$ at $p < 0.05$ rejects the type-agnostic baseline as no-regret. (ii) *Regime-ordering test for (P3):* the per-backbone β -sweep is expected to exhibit a three-regime structure (burn-in-dominated, transition, asymptotic) ordered by per-backbone posterior concentration speed; we test ordering by Kendall’s τ against measured concentration rate from E1. (iii) *Scaling tests for (P1):* E2 and E3 directly fit $|\Theta_i|$ - and n -scaling exponents against the theorem’s $\Omega(|\Theta_i|^n)$ vs $O(n|\Theta_i|)$ separation. Concrete numerical realisations are reported per-experiment below.

4.1 HP-SPGG: Controlled LLM-MAS Benchmark

Setup. A coordinator oversees $n=3$ player LLMs across $K=20$ rounds. Each player’s hidden type $\theta_i \in \Theta_i$ is drawn without replacement from $|\Theta_i|=4$ persona archetypes constructed in the SOTOPIA style (Zhou et al., 2024) and spanning the empirical taxonomy of Fischbacher et al. (2001): altruistic builder, conditional cooperator, strategic hedger, free rider. The state space has $|\mathcal{S}|=5$ scenario framings, joint actions are per-agent contribution levels at $|\mathcal{A}_i|=3$. The joint posterior has $|\Theta|^n=64$ profiles; H-PSMG’s factored representation occupies $4 \cdot 3=12$ entries. TI/RL/PF/ISI/CSR hold by construction; TID is verified empirically (calibration with minimum pairwise Hellinger gap ≥ 0.05).




Figure 1: HP-SPGG cumulative Bayesian regret at $K=20$ across four LLM backbones (5 seeds, calibration c19). Three families: *Bayesian type-aware* (blue), *LLM-coordination* (warm), *type-agnostic* (gray). The H-PSMG family attains the lowest non-oracle regret on every backbone, separating from the best LLM-coordination baseline by 6–12 $\times$ and from the best type-agnostic baseline by over 25 $\times$.

Baselines. We compare H-PSMG⁺ ($\beta=0.25$, fixed across backbones) against three families: *Bayesian type-aware* (H-PSMG, Joint-PSRL, MAP-Type-Greedy); *LLM-coordination* (ECON-BNE (Yang et al., 2025), A-ToM at orders $\{0, 1, 2\}$ with Hedge/FTL aggregation (Xu et al., 2026; Cesa-Bianchi et al., 1997), prompted `llm_belief/llm_greedy`); and *type-agnostic* (PSRL-NoType, IQL, Random).

Cross-model results. Figure 1 shows three clean regret clusters on every backbone. The H-PSMG family attains 0.40 on DeepSeek, 0.64 on GPT-5.4-nano, 0.63 on Kimi, and 0.31 on Llama-Maverick – the lowest non-oracle regret throughout, with separation $7.7\times/6.3\times/11.8\times/10.8\times$ over the best LLM-coordination baseline and 27–35 $\times$ over the best type-agnostic baseline. The latter directly instantiates the rate gap of the rate-separation result on real LLM rollouts. The identity of the best LLM-coordination baseline varies across backbones (`llm_belief` on DeepSeek, A-ToM-0 on GPT-5.4-nano, A-ToM hedge on Kimi, ECON-BNE on Llama-Maverick) – ECON-BNE’s regret on GPT-5.4-nano collapses to 14.88 (vs. 3.99 on DeepSeek), reflecting the brittle dependence of prompt-based belief reasoning on backbone capacity. Closed-form Bayesian belief tracking in H-PSMG is structurally robust.

β -sweep matches Theorem E.4. Figure 2 sweeps β on $\{0, 0.05, 0.1, 0.25, 0.5, 0.75, 1.0, 1.5\}$. Three qualitatively distinct regimes appear, in one-to-one correspondence with the theorem’s predictions. On Llama-Maverick the bonus delivers a




Figure 2: β -sweep of H-PSMG⁺ on HP-SPGG c19 across four backbones. Three regimes predicted by Theorem E.4 appear: *extreme self-decay* (GPT, Kimi); *self-decay edge* (DeepSeek); *burn-in jump* (Llama-Maverick, 3 $\times$ improvement at $\beta=0.25$).




Figure 3: E1: posterior concentration speed predicts the H-PSMG⁺ bonus regime. x -axis: mean $K_{\text{conc}}^{0.9}$ from vanilla H-PSMG ($\beta=0$). y -axis: bonus benefit $\text{Reg}(\text{H-PSMG}) - \text{Reg}(\text{H-PSMG}^+)$ at $\beta=0.25$. Bands are the three regimes of (P3) in §3.3.

sharp 3 $\times$ drop from 0.97 to 0.31 at $\beta=0.25$ – the empirical signature of the exponential burn-in improvement of Theorem E.4(ii). On DeepSeek the regret is flat across $\beta \in [0, 0.75]$ then rises – the high- β boundary of (i). On GPT-5.4-nano and Kimi-K2.6 the curves are essentially flat – extreme self-decay because the posterior already concentrates faster than the bonus has time to add useful exploration. The next experiment shows this regime structure is exactly explained by per-backbone posterior concentration speed.

4.2 Theory-driven Experiments

E1. Posterior concentration speed predicts the regime. Theorem E.4 predicts the bonus contributes nothing once the posterior concentrates, and gives exponential burn-in improvement when concentration is slow. We measure per-backbone concentration time $K_{\text{conc}}^{0.9} = \min\{k : \mu_k^{\Theta, i}(\theta_i^*) \geq 0.9 \forall i\}$ from vanilla H-PSMG ($\beta=0$, independent of the bonus). The four backbones span 3 $\times$: DeepSeek 4.4, GPT 6.6, Kimi 12.0, Llama 14.0 (out of $K=20$). Figure 3 plots concentration time against the observed bonus benefit and recovers the three predicted regimes – the fastest-concentrating backbone




Figure 4: Consolidated empirical evidence for Theorems E.4 and E.3. (a) E2: bonus gap grows from +0.14 at $|\Theta_i|=2$ to +1.08 at $|\Theta_i|=6$ on Llama-Maverick, tracing the predicted $\Omega(|\Theta_i|^n)$ -vs- $O(1)$ envelope of Theorem E.4(ii). (b) E3: bonus benefit grows monotonically with $n \in \{3, 4, 5\}$ while storage savings scale analytically as $|\Theta_i|^n/(n|\Theta_i|)$ ($43\times$ at $n=5$). (c) E5: per-episode cumulative-regret trajectories on HP-SPGG across four backbones; H-PSMG and H-PSMG⁺ plateau at < 0.05 regret per episode within $K \approx 5$ (the $\tilde{O}(\sqrt{K})$ envelope of Theorem E.1) while PSRL-NoType grows linearly (0.50 – 0.93 per episode), realising the rate separation of Theorem E.3 with cumulative gap 11 – $30\times$ at $K=20$. (d) E4: per-backbone burn-in episode $K_{\text{conc}}^{0.9}$ matches the regime ordering of Figure 3. Cross-checks against the analytic kernel (E2-suppl) and against cross-provider rollouts (E3-suppl) appear in Appendix L.

in self-decay edge, the slowest in burn-in jump, the intermediates in a high- β over-exploration boundary regime.

E2–E5 consolidated. Figure 4 reports the remaining theory-driven experiments. Panel (a) varies $|\Theta_i| \in \{2, 3, 4, 5, 6\}$ at $n=3$ on Llama-Maverick and realises Theorem E.4(ii)’s two-regime prediction: vanilla H-PSMG tracks the $\Omega(|\Theta_i|^n)$ envelope while H-PSMG⁺ stays $O(1)$. Panel (b) varies $n \in \{3, 4, 5\}$ at fixed $|\Theta_i|=4$, with the H-PSMG⁺ benefit growing monotonically and the factored-posterior storage advantage growing as $|\Theta_i|^n/(n|\Theta_i|)$. Panel (c) (E5) reports per-episode trajectories across all four backbones: H-PSMG⁺ and H-PSMG plateau at < 0.05 regret per episode within $K \approx 5$ on every backbone – the $\tilde{O}(\sqrt{K})$ envelope of Theorem E.1 realised, with the empirical regime dominated by gap-dependent identifiability so the burn-in term is paid once. PSRL-NoType accumulates 0.50 – 0.93 per episode and grows linearly

across all backbones; the $K=20$ cumulative gap of 11 – $30\times$ is the empirical signature of Theorem E.3’s rate separation on the LLM substrate. Panel (d) reports the per-backbone burn-in episode $K_{\text{conc}}^{0.9}$ in trajectory space, recovering the same regime ordering as Figure 3. Two further cross-checks – (i) H-PSMG⁺ matches a privileged analytic-kernel reference within 0.05 regret across all four backbones, and (ii) cross-provider rollouts that mix two backbones across the $n=3$ slots retain the same regret ordering – are deferred to Appendix L.

4.3 Concordia Public Benchmarks

To rule out benchmark bias we evaluate on the public Concordia substrates (Vezhnevets et al., 2023; Smith et al., 2025), independently designed by DeepMind. For each substrate we map agent disposition to θ_i , exogenous state to s , and per-agent decisions to a_i ; Concordia provides analytic reward via the substrate’s coordination payoff component, eliminating the LLM judge. We compare H-PSMG⁺ (joint proxy and per-agent MAP variants) against ECON-BNE-mech and A-ToM-1-mech adapted to the Concordia interface.

Figure 5 reports focal-score means across 7 Pub Coordination configurations spanning four village variants and three closure regimes, and 9 Hagglng sub-configurations (fruitville and vegbrooke variants, single and multi-item). On Pub Coordination, H-PSMG⁺ (joint proxy) is within 1% of `oracle_joint` on every mechanistic config; ECON-BNE-mech trails by 5–15%, A-ToM-1-mech by 5–30%. On Hagglng, H-PSMG⁺ wins or ties on all 9 configs and *strictly exceeds* `oracle_joint` on three (vegbrooke_stubborn +1.43, two fruitville_gullible +0.40, +0.47): `oracle_joint` maximises joint payoff while H-PSMG⁺ maximises focal-agent payoff. On vegbrooke_strange, A-ToM-1-mech collapses to the substrate floor at -10.0 while H-PSMG⁺ holds at 0 – a robustness gap from closed-form Bayesian belief tracking.

4.4 SOTOPIA-Hard: Partial Transfer Beyond Theoretical Preconditions

SOTOPIA-Hard (Zhou et al., 2024) is a dyadic ($n=2$) social-negotiation benchmark with open-ended natural-language actions and a persona structure that violates the prior-factorisation assumption PF. The substrate is structurally out of scope for HT-MG-BC on three axes: it is dyadic rather than $n \geq 3$, the action space is unbounded text rather than tabular contribution levels, and the persona prior is not factored. We use SOTOPIA-Hard to probe whether the H-PSMG⁺ framework remains useful where its theoretical preconditions are out of scope. A substrate-adapted variant extracts approximate per-agent marginals from a Concordia-style mechanistic surrogate over a discretised intent space, then applies the same factored-posterior




Figure 5: Concordia results. *Left*: Pub Coordination, 9 axes (7 village configs at $s=30$, one at $s=100$, one DeepSeek-V3.2 live LLM rollout at $s=5$). *Right*: Hagglng, 9 axes spanning 5 single-item and 4 multi-item configs at $s=30$. Each axis is per-axis min-max normalised so the best of four algorithms reaches the outer ring; oracle is dashed and may sit inside the ring because H-PSMG⁺ can strictly exceed oracle on focal score. H-PSMG⁺ (joint proxy) fills nearly the full Pub Coordination polygon and wins or ties on every Hagglng axis.

sampling and type-discrimination bonus as in HP-SPGG, in Sotopia self-play.

H-PSMG⁺ transfers partially (Figure 6).

Over 70 episodes (14 scenarios $\times$ 5 seeds) per backbone, H-PSMG⁺ remains within 0.1 of the best alternative on every backbone (gap -0.019 on GPT-5.4-nano, up to -0.099 on Kimi-K2.6; Figure 6(A)). On the largest scenario family *craigslst_bargains* (16 episode batches), H-PSMG⁺ leads with the highest win count (5 wins, Figure 6(B)) ahead of every other baseline. It also leads on *sleep_arrangement* and *donate_funds* (2/4 wins each) and ties on *revenge_plot*. On GPT-5.4-nano the focal-agent (agent₁) score is the highest of any baseline (3.122 vs. next-best *llm_belief* at 3.118). On the remaining backbones, H-PSMG⁺ falls into a tight cluster of strong LLM-coordination baselines within the same 0.1-score envelope; the strict dominance visible in HP-SPGG and Concordia does not persist out of scope. This is exactly the boundary that the prior-factorisation lower bound (Theorem E.5) predicts: the exponential separation requires PF. The framework therefore degrades gracefully rather than catastrophically and remains in win-count parity with the strongest baselines, with full per-(backbone, scenario) breakdown deferred to Appendix M.

5 Discussion, Limitations, and Conclusion

Summary. We have introduced H-PSMG and H-PSMG⁺, an algorithm family for centralised multi-agent LLM coordination under unknown personas and dynamics. The factored posterior is exact under three structural conditions and yields exponential storage savings in n together with a Bayesian regret bound that is exponentially tighter than the prior-factorisation lower bound. The type-discrimination bonus reduces burn-in regret by an $\Omega(|\Theta_i|^n)$ factor on a deep-trap family. Across three LLM benchmarks of increasing wildness – HP-SPGG, Concordia, SOTOPIA-Hard – the algorithm family beats the strongest published LLM-coordination baselines by 6–12 $\times$ where the theoretical preconditions hold (HP-SPGG), wins or ties oracle on 16/16 Concordia configurations, and remains within 0.1 of the best alternative on every backbone in the out-of-scope dyadic-negotiation setting (SOTOPIA-Hard) where it leads the largest scenario family while losing the strict dominance the theory guarantees only under prior factorisation – graceful degradation rather than catastrophic failure as the preconditions weaken.

Limitations. (i) The regret guarantees assume a tabular state-action space and a discrete persona archetype set; the function-approximation generalisation FA-H-PSMG⁺ (Appendix K) extends to con-




Figure 6: SOTOPIA-Hard results. (A) Aggregate mean focal score per (backbone $\times$ baseline), 70 episodes per cell. H-PSMG⁺ (deep navy, bold outline) stays within 0.1 of the leader on every backbone (numerical scores marked); the five baselines cluster in a 0.12-score envelope, indicating that out of scope no algorithm achieves a clear separation – consistent with the prior-factorisation lower bound. (B) Wins per scenario family across the 56 (backbone, scenario codename) pairs. H-PSMG⁺ leads on the largest family *craigslist_bargains* (5/16) and on *sleep_arrangement* and *donate_funds* (2/4 each), ties on *revenge_plot*, and is competitive (1 win) on *divide_items* and *join_trip*.

tinuous spaces but inherits function-approximation error. (ii) The CCE oracle is assumed deterministic and measurable (CSR); randomised CCE-tie-breaking would alter the burn-in constants. (iii) Personas are sampled once per episode and held fixed; non-stationary personas are out of scope. (iv) The persona-likelihood scoring requires next-token-probability access from the LLM backbone, which all four backbones we tested support; black-box-only LLMs would require a wrapper. (v) The bonus weight β is held fixed at 0.25 across all substrates and backbones; per-substrate adaptation could further improve burn-in performance but at the cost of substrate-specific tuning.

Future work. Decentralised HT-MG-BC (each agent maintains its own posterior over the others), beyond-TI dynamics (persona-coupled transitions), and adaptive- β schedules informed by per-episode posterior concentration estimates are natural next steps. The function-approximation generalisation (Appendix K) is fully worked out and ready for empirical study at scale.

References

- Shipra Agrawal and Randy Jia. 2017. Posterior sampling for reinforcement learning: worst-case regret bounds. *Advances in Neural Information Processing Systems*, 30.
- Stefano V Albrecht, Jacob W Crandall, and Subramanian Ramamoorthy. 2019. On convergence and optimality of best-response learning with policy types in multiagent systems. *Conference on Uncertainty in Artificial Intelligence (UAI) / arXiv:1907.06995*.
- Stefano V Albrecht and Subramanian Ramamoorthy. 2013. A game-theoretic model and best-response learning method for ad hoc coordination in multiagent systems. In *AAMAS*.
- Dilip Arumugam and Thomas L. Griffiths. 2025. Toward efficient exploration by large language model agents. In *International Conference on Machine Learning (ICML)*.
- Dilip Arumugam and Thomas L Griffiths. 2026. Toward efficient exploration by large language model agents. In *International Conference on Learning Representations*.
- Mohammad Gheshlaghi Azar, Ian Osband, and Rémi Munos. 2017. Minimax regret bounds for reinforcement learning. In *International Conference on Machine Learning*, pages 263–272.
- Yu Bai and Chi Jin. 2020. Provable self-play algorithms for competitive reinforcement learning. In *International Conference on Machine Learning*.
- Nicolò Cesa-Bianchi, Yoav Freund, David Haussler, David P Helmbold, Robert E Schapire, and Manfred K Warmuth. 1997. How to use expert advice. *Journal of the ACM*, 44(3):427–485.
- Thomas M. Cover and Joy A. Thomas. 2006. *Elements of Information Theory*, 2nd edition. Wiley-Interscience.
- Qiwen Cui, Kaiqing Yang, and Simon S Du. 2025. Incentivize without bonus: Provably efficient model-based online multi-agent rl for markov games. *arXiv preprint arXiv:2502.09780*.
- Le Cong Dinh, Yaodong Yang, Stephen McAleer, Zheng Tian, Nicolas Perez-Nieves, Oliver Slumbers, David Henry Mguni, Haitham Bou-Ammar, and Jun Wang. 2022. Online double oracle.

- In *Transactions on Machine Learning Research (TMLR)*.
- Omar Darwiche Domingues, Pierre Ménard, Emilie Kaufmann, and Michal Valko. 2021. Episodic reinforcement learning in finite MDPs: Minimax lower bounds revisited. In *Algorithmic Learning Theory (ALT)*, pages 578–598.
- Liad Erez, Tal Lancewicki, Uri Sherman, Tomer Koren, and Yishay Mansour. 2023. Regret minimization and convergence to equilibria in general-sum markov games. In *International Conference on Machine Learning*.
- Urs Fischbacher, Simon Gächter, and Ernst Fehr. 2001. Are people conditionally cooperative? evidence from a public goods experiment. *Economics letters*, 71(3):397–404.
- Yarin Gal and Zoubin Ghahramani. 2016. Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. In *ICML*.
- Subhashis Ghosal and Aad van der Vaart. 2017. *Fundamentals of Nonparametric Bayesian Inference*. Cambridge University Press.
- John C Harsanyi. 1967,1968. Games with incomplete information played by bayesian players, i-iii. *Management science*, 14(3,5,7):159–182, 320–334, 486–502.
- Hao-Lun Hsu, Weixin Wang, Miroslav Pajic, and Pan Xu. 2024. Randomized exploration in cooperative multi-agent reinforcement learning. In *Advances in Neural Information Processing Systems*.
- Mehdi Jafarnia-Jahromi, Rahul Jain, and Ashutosh Nayyar. 2021. Learning zero-sum stochastic games with posterior sampling. *arXiv preprint arXiv:2109.03396*.
- Thomas Jaksch, Ronald Ortner, and Peter Auer. 2010. Near-optimal regret bounds for reinforcement learning. *Journal of Machine Learning Research*, 11:1563–1600.
- Chi Jin, Qinghua Liu, Yuanhao Wang, and Tiancheng Yu. 2021. V-learning: A simple, efficient, decentralized algorithm for multi-agent rl. In *Advances in Neural Information Processing Systems*.
- Michael Kearns and Satinder Singh. 2002. Near-optimal reinforcement learning in polynomial time. *Machine Learning*, 49:209–232.
- Jeongyeol Kwon, Yonathan Efroni, Constantine Caramanis, and Shie Mannor. 2021. RL for latent mdp: Regret guarantees and a lower bound. In *Advances in Neural Information Processing Systems*.
- Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. 2017. Simple and scalable predictive uncertainty estimation using deep ensembles. In *NeurIPS*.
- Marc Lanctot, Vinicius Zambaldi, Audrunas Gruslys, Angeliki Lazaridou, Karl Tuyls, Julien Pérolat, David Silver, and Thore Graepel. 2017. A unified game-theoretic approach to multiagent reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Tor Lattimore and Csaba Szepesvári. 2020. *Bandit Algorithms*. Cambridge University Press.
- Joel Z Leibo, Vinicius Zambaldi, Marc Lanctot, Janusz Marecki, and Thore Graepel. 2017. Multi-agent reinforcement learning in sequential social dilemmas. In *Proceedings of AAMAS*.
- Michael L Littman. 1994. Markov games as a framework for multi-agent reinforcement learning. In *Machine learning proceedings 1994*, pages 157–163.
- Xiuyuan Lu and Benjamin Van Roy. 2019. Information-theoretic confidence bounds for reinforcement learning. *Advances in Neural Information Processing Systems*, 32.
- Weichao Mao and Tamer Başar. 2023. Provably efficient reinforcement learning in decentralized general-sum markov games. *Dynamic Games and Applications*.
- Luke Marris, Paul Muller, Marc Lanctot, Karl Tuyls, and Thore Graepel. 2021. Multi-agent training beyond zero-sum with correlated equilibrium meta-solvers. In *ICML*.
- Stephen McAleer, John Lanier, Pierre Baldi, and Roy Fox. 2021. XDO: A double oracle algorithm for extensive-form games. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Stephen McAleer, John Lanier, Roy Fox, and Pierre Baldi. 2020. Pipeline PSRO: A scalable approach for finding approximate Nash equilibria in large games. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Stephen McAleer, Kevin Wang, John Lanier, Marc Lanctot, Pierre Baldi, Tuomas Sandholm, and Roy Fox. 2022. Anytime PSRO for two-player zero-sum games. In *AAAI Workshop on Reinforcement Learning in Games*.
- Paul Muller, Shayegan Omidshafiei, Mark Rowland, Karl Tuyls, Julien Pérolat, Siqi Liu, Daniel Hennes, Luke Marris, Marc Lanctot, Edward Hughes, and 1 others. 2020. A generalized training approach for multiagent learning. In *ICLR*.
- Mirco Mutti, Riccardo De Meo, and Matteo Pirotta. 2024. Exploiting causal graph priors with posterior sampling for reinforcement learning. *arXiv preprint arXiv:2310.07518*.
- Ian Osband, Daniel Russo, and Benjamin Van Roy. 2013. (more) efficient reinforcement learning via posterior sampling. In *Advances in Neural Information Processing Systems*.

- Ian Osband and Benjamin Van Roy. 2014. Model-based reinforcement learning and the eluder dimension. In *Advances in Neural Information Processing Systems*.
- Ian Osband and Benjamin Van Roy. 2017. Why is posterior sampling better than optimism for reinforcement learning? In *International Conference on Machine Learning*, pages 2701–2710. PMLR.
- Yi Ouyang, Mukul Gagrani, Ashutosh Nayyar, and Rahul Jain. 2017. Learning unknown markov decision processes: A thompson sampling approach. In *Advances in Neural Information Processing Systems*.
- Chanwoo Park, Xiangyu Liu, Asuman Ozdaglar, and Kaiqing Zhang. 2025. Do LLM agents have regret? A case study in online learning and games. In *International Conference on Learning Representations (ICLR)*.
- Julien Perolat, Joel Z Leibo, Vinicius Zambaldi, Charles Beattie, Karl Tuyls, and Thore Graepel. 2017. A multi-agent reinforcement learning model of common-pool resource appropriation. In *Advances in Neural Information Processing Systems*.
- Daniel Russo and Benjamin Van Roy. 2014. Learning to optimize via posterior sampling. *Mathematics of Operations Research*, 39(4):1221–1243.
- Daniel Russo and Benjamin Van Roy. 2016. An information-theoretic analysis of thompson sampling. *Journal of Machine Learning Research*, 17(1):2442–2471.
- Lloyd S Shapley. 1953. Stochastic games. *Proceedings of the National Academy of Sciences*, 39(10):1095–1100.
- Adam Smith and 1 others. 2025. Evaluating generalization capabilities of LLM-based agents in mixed-motive scenarios using Concordia. In *NeurIPS Datasets and Benchmarks Track*. Concordia Contest 2024 results paper.
- Ziang Song, Song Mei, and Yu Bai. 2022. When can we learn general-sum markov games with a large number of players sample-efficiently? In *International Conference on Learning Representations*.
- Malcolm Strens. 2000. A bayesian framework for reinforcement learning. In *ICML*, volume 2000, pages 943–950.
- Yuchen Tian, Tiancheng Yu, Hao Hu, and Chi Jin. 2024. Learning in markov games with adaptive adversaries: Policy regret, fundamental barriers, and efficient algorithms. *arXiv preprint arXiv:2411.00707*.
- Alexander Sasha Vezhnevets, John P Agapiou, Avia Aharon, Ron Ziv, Jayd Matyas, Edgar A Duéñez-Guzmán, William A Cunningham, Simon Osindero, Danny Karmon, and Joel Z Leibo. 2023. Generative agent-based modeling with actions grounded in physical, social, or digital space using Concordia. *arXiv preprint arXiv:2312.03664*.
- Tsachy Weissman, Erik Ordentlich, Gadiel Seroussi, Sergio Verdú, and Marcelo J Weinberger. 2003. Inequalities for the L1 deviation of the empirical distribution. In *Hewlett-Packard Labs, Tech. Rep.*
- Max Welling and Yee Whye Teh. 2011. Bayesian learning via stochastic gradient langevin dynamics. In *ICML*.
- Wei Xiong, Han Zhong, Chengshuai Shi, Cong Shen, and Tong Zhang. 2022. A self-play posterior sampling algorithm for zero-sum markov games. In *International Conference on Machine Learning*.
- Ziyao Xu and 1 others. 2026. Adaptive theory of mind for LLM-based multi-agent coordination. In *AAAI Conference on Artificial Intelligence*.
- Xie Yang and 1 others. 2025. From debate to equilibrium: Belief-driven multi-agent LLM reasoning via Bayesian Nash equilibrium. In *International Conference on Machine Learning (ICML)*.
- Jian Yao, Weiming Liu, Haobo Fu, Yaodong Yang, Stephen McAleer, Qiang Fu, and Wei Yang. 2023. Policy space diversity for non-transitive games. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Chang-Wei Yueh, Andy Zhao, Ashutosh Nayyar, and Rahul Jain. 2026. Bayesian learning in episodic zero-sum games. *arXiv preprint arXiv:2603.20604*.
- Kaiqing Zhang, Zhuoran Yang, and Tamer Başar. 2021. Multi-agent reinforcement learning: A selective overview of theories and algorithms. *Handbook of reinforcement learning and control*, pages 321–384.
- Ming Zhou, Jingxiao Chen, Ying Wen, Weinan Zhang, Yaodong Yang, Yong Hong, and Jun Wang. 2022. Efficient policy space response oracles. *arXiv preprint arXiv:2202.00633*.
- Xuhui Zhou, Hao Zhu, Leena Mathur, Ruohong Zhang, Haofei Yu, Zhengyang Qi, Louis-Philippe Morency, Yonatan Bisk, Daniel Fried, Graham Neubig, and Maarten Sap. 2024. SOTOPIA: Interactive evaluation for social intelligence in language agents. In *International Conference on Learning Representations (ICLR)*.

A HT-MG-BC Formal Setting

Heterogeneous-Type Markov Game. An HT-MG is a tuple $\mathcal{G} = \langle \mathcal{S}, \{\mathcal{A}_i\}, P^*, \{\Theta_i\}, \{r_i\}, P_0, H \rangle$ with finite state space $|\mathcal{S}| = S$, joint action space $\mathcal{A} = \prod_i \mathcal{A}_i$ of size A_{joint} , type space $\Theta = \prod_i \Theta_i$, transition kernel $P^*: \mathcal{S} \times \mathcal{A} \times \Theta \rightarrow \Delta(\mathcal{S})$, per-agent reward $r_i: \mathcal{S} \times \mathcal{A} \times \Theta_i \rightarrow [0, 1]$, joint prior P_0 , and horizon H . At the start of each episode a type

profile $\theta^* \in \Theta$ is sampled once and held fixed; P^* is fixed across episodes but unknown. A centralised learner observes the full trajectory τ_k and updates a posterior over (P^*, θ^*) .

Structural assumptions. The decoupling theorem rests on three: **(TI)**, Assumption A.1 the transition kernel does not depend on the type profile; **(RL)**, Assumption A.2 each r_i depends only on θ_i (other agents' actions a_{-i} may still influence reward); **(PF)**, Assumption A.3 the prior factorises as $P_0(P, \theta) = P_0^P(P) \prod_i P_0^{\Theta, i}(\theta_i)$. The regret analysis additionally uses type-identifiability **(TID)**, Assumption A.4: per-step Hellinger gap $\geq \rho$, initial-state independence **(ISI)**, Assumption A.5, and a deterministic CCE selection rule **(CSR)**, Assumption A.6. Both TI and RL are individually necessary: violating either breaks the likelihood factorisation, as shown by two counter-examples in Appendix D.

Assumption A.1 (TI). $P^*(s' | s, a, \theta) = P^*(s' | s, a)$.

Assumption A.2 (RL). $r_i(s, a, \theta) = r_i(s, a, \theta_i)$.

Assumption A.3 (PF). $P_0(P, \theta) = P_0^P(P) \prod_i P_0^{\Theta, i}(\theta_i)$.

Assumption A.4 (TID). $\inf_{s, a} D_H^2(P(\cdot | s, a, \theta_i), P(\cdot | s, a, \theta'_i)) \geq \rho > 0$ for all $\theta_i \neq \theta'_i$.

Assumption A.5 (ISI). $s_1^k \sim \beta$ fixed and independent of (P^*, θ^*) .

Assumption A.6 (CSR). CCE($\cdot, \cdot$) is a measurable deterministic mapping.

Performance criterion. Let $\pi^{\text{CCE}}(P, \theta)$ denote a CCE policy of the MG induced by (P, θ) . The type-conditional Bayesian regret of agent i over K episodes is

$$\text{Reg}_i^B(K) := \mathbb{E}_{(P^*, \theta^*) \sim P_0} \sum_{k=1}^K \left(V_i^{\pi^{\text{CCE}}(P^*, \theta^*)}(s_1^k; \theta_i^*) - V_i^{\pi_k}(s_1^k; \theta_i^*) \right). \quad (2)$$

B Posterior Decoupling Theorem

This section states and proves the main structural result.

B.1 Likelihood Factorization

We first derive the conditional likelihood of an observed trajectory given the latent world. Under H-PSMG, the policy π_k at episode k is a function only of the history $\mathcal{H}_{k-1}$ and the algorithmic posterior samples $(\hat{P}_k, \hat{\theta}_k)$, which are themselves measurable with respect to $\mathcal{H}_{k-1}$. Crucially, π_k does not depend on the true latent (P^*, θ^*) .

Lemma B.1 (Joint Likelihood Factorization). Under Assumptions A.1 and A.2, the trajectory likelihood

at episode k factorizes as

$$P(\tau_k | P, \theta, \mathcal{H}_{k-1}) = c_k(\tau_k) \cdot L_P(\tau_k; P) \cdot \prod_{i \in [n]} L_{\Theta, i}(\tau_k; \theta_i) \quad (3)$$

where

$$c_k(\tau_k) := \mathbb{E}_{\pi_k | \mathcal{H}_{k-1}} \left[\prod_{h=1}^H \pi_k(a_h^k | s_h^k) \right], \quad (4)$$

$$L_P(\tau_k; P) := \prod_{h=1}^H P(s_{h+1}^k | s_h^k, a_h^k), \quad (5)$$

$$L_{\Theta, i}(\tau_k; \theta_i) := \prod_{h=1}^H P(r_h^{k, i} | s_h^k, a_h^k, \theta_i). \quad (6)$$

The constant $c_k(\tau_k)$ depends on τ_k and $\mathcal{H}_{k-1}$ but not on (P, θ) .

Proof. Conditional on π_k , the trajectory likelihood unrolls along the time axis as

$$P(\tau_k | \pi_k, P, \theta) = \prod_{h=1}^H \pi_k(a_h^k | s_h^k) \cdot P(s_{h+1}^k | s_h^k, a_h^k, \theta) \cdot \prod_{i \in [n]} P(r_h^{k, i} | s_h^k, a_h^k, \theta_i) \quad (7)$$

Apply Assumption A.1 to drop θ from the transition factor, and Assumption A.2 to localize the reward factor to θ_i :

$$P(\tau_k | \pi_k, P, \theta) = \prod_{h=1}^H \pi_k(a_h^k | s_h^k) \cdot \prod_{h=1}^H P(s_{h+1}^k | s_h^k, a_h^k) \cdot \prod_{h=1}^H \prod_{i \in [n]} P(r_h^{k, i} | s_h^k, a_h^k, \theta_i) \quad (8)$$

Marginalize over π_k with respect to $P(\pi_k | \mathcal{H}_{k-1})$, which is independent of (P, θ) :

$$P(\tau_k | P, \theta, \mathcal{H}_{k-1}) = \int P(\pi_k | \mathcal{H}_{k-1}) \cdot P(\tau_k | \pi_k, P, \theta) d\pi_k \quad (9)$$

$$= \mathbb{E}_{\pi_k | \mathcal{H}_{k-1}} \left[\prod_{h=1}^H \pi_k(a_h^k | s_h^k) \right] \cdot L_P(\tau_k; P) \cdot \prod_{i \in [n]} L_{\Theta, i}(\tau_k; \theta_i) \quad (10)$$

The factors L_P and $L_{\Theta, i}$ depend respectively on P alone and θ_i alone, and the leading expectation depends on τ_k and $\mathcal{H}_{k-1}$ alone. $\square$

B.2 Main Theorem

Theorem B.2 (Posterior Decoupling). Under Assumptions A.1, A.2, and A.3, the posterior distribution over the latent world factorizes at every episode:

$$P(P, \theta | \mathcal{H}_k) = \mu_{k+1}^P(P) \cdot \prod_{i \in [n]} \mu_{k+1}^{\Theta, i}(\theta_i), \quad \forall k \in \{0, 1, \dots, K\} \quad (11)$$

where the marginals admit the recursive updates

$$\mu_{k+1}^P(P) \propto \mu_k^P(P) \cdot L_P(\tau_k; P), \quad \mu_{k+1}^{\Theta, i}(\theta_i) \propto \mu_k^{\Theta, i}(\theta_i) \cdot L_{\Theta, i}(\tau_k; \theta_i) \quad (12)$$

with initial conditions $\mu_1^P = P_0^P$ and $\mu_1^{\Theta, i} = P_0^{\Theta, i}$.

Proof. We proceed by induction on k .

Base case ($k = 0$). By Assumption A.3, $P(P, \theta | \mathcal{H}_0) = P_0(P, \theta) = P_0^P(P) \prod_{i \in [n]} P_0^{\Theta, i}(\theta_i)$, which has the required factored form.

Inductive step. Suppose $P(P, \theta | \mathcal{H}_{k-1}) = \mu_k^P(P) \prod_i \mu_k^{\Theta, i}(\theta_i)$. By Bayes' rule,

$$P(P, \theta | \mathcal{H}_k) \propto P(P, \theta | \mathcal{H}_{k-1}) \cdot P(\tau_k | P, \theta, \mathcal{H}_{k-1}). \quad (13)$$

Apply Lemma B.1:

$$P(P, \theta | \mathcal{H}_k) \propto \mu_k^P(P) \prod_i \mu_k^{\Theta, i}(\theta_i) \cdot c_k(\tau_k) \cdot L_P(\tau_k; P). \quad (14)$$

$$\propto [\mu_k^P(P) L_P(\tau_k; P)] \cdot \prod_i [\mu_k^{\Theta, i}(\theta_i) L_{\Theta, i}(\tau_k; \theta_i)]. \quad (15)$$

The factor $c_k(\tau_k)$ does not depend on (P, θ) and is absorbed into the normalization constant. The remaining product factorizes into a P -only term and n separate θ_i -only terms. Normalizing each factor independently yields the recursive forms stated in the theorem and establishes the desired posterior factorization. $\square$

Corollary B.3 (Storage and Update Complexity). *Maintaining the joint posterior $P(P, \theta | \mathcal{H}_k)$ under H-PSMG requires $O(\dim(P_0^P) + n|\Theta|)$ storage, as opposed to the naive $O(\dim(P_0^P) \cdot |\Theta|^n)$ that a non-factored posterior would require. Each Bayesian update costs $O(H)$ likelihood evaluations per module.*

Proposition B.4 (Trajectory-Level Equivalence with Joint-PSRL). *Consider an idealised Joint-PSRL algorithm that maintains the joint posterior μ_k^Θ on $\Theta = \prod_i \Theta_i$ and at each episode samples $\hat{\theta}_k \sim \mu_k^\Theta$. Under the assumptions of Theorem B.2 and coupled RNG (formalised below), the trajectories produced by H-PSMG and Joint-PSRL are pathwise identical: $\tau_k^{\text{H-PSMG}} = \tau_k^{\text{Joint-PSRL}}$ for every $k \geq 1$ almost surely. In particular the cumulative regrets agree exactly:*

$$\text{Reg}_i^B[\text{H-PSMG}](K) = \text{Reg}_i^B[\text{Joint-PSRL}](K) \quad \forall i, K \in \mathbb{N}. \quad (16)$$

The coupling is constructed as follows. At episode k , both algorithms draw a single shared seed $u_k = (u_k^P, u_k^{\Theta, 1}, \dots, u_k^{\Theta, n}) \sim \text{Unif}([0, 1]^{n+1})$. H-PSMG uses u_k^P to sample $\hat{P}_k = F_{\mu_k^P}^{-1}(u_k^P)$ and $u_k^{\Theta, i}$ to sample $\hat{\theta}_i^k = F_{\mu_k^{\Theta, i}}^{-1}(u_k^{\Theta, i})$ via the inverse-CDF transform. Joint-PSRL uses u_k^P identically for $\hat{P}_k$ and the joint draw $\hat{\theta}_k = F_{\mu_k^\Theta}^{-1}(u_k^{\Theta, 1}, \dots, u_k^{\Theta, n})$ via the per-coordinate inverse-CDF of the product factorization.

Proof. By Theorem B.2, $\mu_k^\Theta(\theta_1, \dots, \theta_n) = \prod_i \mu_k^{\Theta, i}(\theta_i)$ at every k . The inverse-CDF of a

product distribution applied coordinate-wise yields independent samples from the marginals. Hence the two algorithms produce identical $\hat{P}_k$ and identical $(\hat{\theta}_1^k, \dots, \hat{\theta}_n^k)$, hence identical CCE policies $\pi_k = \text{CCE}(\hat{P}_k, \hat{\theta}_k)$ (Assumption A.6: deterministic oracle), hence identical trajectories τ_k under shared environment randomness. Identical trajectories imply identical posteriors at $k+1$ by Bayes' rule, and the argument repeats inductively. $\square$

Proposition B.4 formalises that H-PSMG and Joint-PSRL produce bit-identical trajectories under matched RNG on any HT-MG-BC instance. The proposition shows that H-PSMG is not merely a heuristic approximation of Joint-PSRL but its exact statistical equivalent at $O(n|\Theta_i|)$ rather than $O(|\Theta|^n)$ memory. The qualitative implication is that all regret guarantees and lower bounds derived for H-PSMG transfer verbatim to Joint-PSRL on HT-MG-BC instances.

Remark B.5. *The decoupling is more than a notational convenience. It separates two distinct sources of uncertainty into modular inference problems. Model uncertainty is captured by μ^P , which behaves exactly as in single-agent PSRL and inherits all known concentration tools (Agrawal and Jia, 2017). Type uncertainty is captured by $\mu^{\Theta, i}$, which is a categorical Bayesian estimation problem amenable to information-theoretic analysis (Russo and Van Roy, 2014). This separation is the basis of the regret decomposition outlined in Section E.*

C Numerical Illustration: Two-Agent Bayesian Prisoner's Dilemma

To make the decoupling concrete we work through a minimal instance numerically. The example is a stateless two-agent two-type matrix game played for K rounds.

C.1 Setup

Let $|\mathcal{S}| = 1$, $\mathcal{A}_1 = \mathcal{A}_2 = \{C, D\}$, and $\Theta_i = \{P, S\}$ for $i \in \{1, 2\}$, where P denotes pro-social and S denotes selfish. The reward of agent i as a function of its type is

$$r_i^S(a_i, a_{-i}) = \mathbb{1}[a_i = D], \quad r_i^P(a_i, a_{-i}) = \mathbb{1}[a_i = C \wedge a_{-i} = C]. \quad (17)$$

A selfish agent prefers defection. A pro-social agent rewards mutual cooperation. The horizon is $H = 1$ and there is no transition uncertainty, so the model component μ^P is trivial and the example isolates the type-inference component.

The prior is uniform: $\mu_1^{\Theta, 1}(P) = \mu_1^{\Theta, 1}(S) = 0.5$ and similarly for agent 2, so $\mu_1^\Theta(\theta) = 0.25$ for each $\theta \in \Theta$. The true type profile is $\theta^* = (P, S)$.

C.2 Episode 1

The algorithm draws $\hat{\theta}_1 \sim \mu_1^\Theta$ uniformly. Suppose the realization is $\hat{\theta}_1 = (P, P)$. The CCE of the all-pro-social game is the pure strategy (C, C) . The algorithm executes (C, C) , and rewards are realized in the true world: $r_1 = r_1^P(C, C) = 1$ and $r_2 = r_2^S(C, \cdot) = 0$.

The reward likelihoods are

$$L_{\Theta,1}(\tau_1; \theta_1) = P(r_1 = 1 \mid a = (C, C), \theta_1) = \begin{cases} 1 & \theta_1 = P \\ 0 & \theta_1 = S \end{cases} \quad (18)$$

$$L_{\Theta,2}(\tau_1; \theta_2) = P(r_2 = 0 \mid a = (C, C), \theta_2) = \begin{cases} 0 & \theta_2 = P \\ 1 & \theta_2 = S \end{cases} \quad (19)$$

Table 1 tabulates the posterior update.

Table 1: Posterior update after episode 1. Joint posterior is computed by Bayes' rule, then verified against the product of marginals.

θ_1	θ_2	$\mu_1^\Theta(\theta)$	$L_{\Theta,1}$	$L_{\Theta,2}$	Unnormalized posterior
P	P	0.25	1	0	0
P	S	0.25	1	1	0.25
S	P	0.25	0	0	0
S	S	0.25	0	1	0
Normalized					
P	S				1.00
others					0.00

The joint posterior puts probability 1 on the true type profile (P, S) . We verify the decoupling by computing the marginals:

$$\mu_2^{\Theta,1}(P) = \mu_1^{\Theta,1}(P) \cdot L_{\Theta,1}(\tau_1; P) / Z_1 = 0.5 \cdot 1 / 0.5 = 1, \quad (20)$$

$$\mu_2^{\Theta,1}(S) = 0, \quad (21)$$

$$\mu_2^{\Theta,2}(P) = \mu_1^{\Theta,2}(P) \cdot L_{\Theta,2}(\tau_1; P) / Z_2 = 0.5 \cdot 0 / 0.5 = 0, \quad (22)$$

$$\mu_2^{\Theta,2}(S) = 1. \quad (23)$$

The product $\mu_2^{\Theta,1} \cdot \mu_2^{\Theta,2}$ equals the unnormalized joint $\delta_{(P,S)}$, confirming the decoupling in Theorem B.2.

C.3 Noisy Observation Variant

To show that decoupling is not an artifact of deterministic rewards, suppose now the reward signal flips with probability $\epsilon = 0.1$. The likelihoods become

$$L_{\Theta,1}(\tau_1; \theta_1) = \begin{cases} 0.9 & \theta_1 = P \\ 0.1 & \theta_1 = S \end{cases}, \quad L_{\Theta,2}(\tau_1; \theta_2) = \begin{cases} 0.1 & \theta_2 = P \\ 0.9 & \theta_2 = S \end{cases} \quad (24)$$

The unnormalized joint posterior is the outer product $\mu_1^\Theta \cdot L_{\Theta,1} \cdot L_{\Theta,2}$ which still factorizes as $\mu_2^{\Theta,1} \cdot \mu_2^{\Theta,2}$ with $\mu_2^{\Theta,1}(P) = 0.9, \mu_2^{\Theta,1}(S) = 0.1$ and $\mu_2^{\Theta,2}(S) = 0.9, \mu_2^{\Theta,2}(P) = 0.1$. Concentration is partial after one observation but the factorization is exact.

D Necessity of the Structural Assumptions

We show that both Assumptions A.1 and A.2 are necessary in the sense that violation of either causes the decoupling to break.

D.1 Counter-Example I: Type-Dependent Transitions

Suppose Assumption A.1 fails. Concretely, consider a two-state HT-MG where $P(s' \mid s, a, \theta)$ varies with θ . The trajectory likelihood becomes

$$P(\tau_k \mid P, \theta, \mathcal{H}_{k-1}) = c_k(\tau_k) \cdot \prod_{h=1}^H P(s_{h+1}^k \mid s_h^k, a_h^k, \theta) \cdot \prod_{h,i} P(r_h^{k,i} \mid s_h^k, a_h^k, \theta) \quad (25)$$

The transition factor is now jointly a function of P and θ and does not split. Posterior $P(P, \theta \mid \mathcal{H}_k)$ is no longer factorable into a P -marginal and a θ -marginal. Empirically, observing a transition $(s, a) \rightarrow s'$ simultaneously updates beliefs about both the kernel and the types, with no way to attribute the change to one source alone.

D.2 Counter-Example II: Non-Local Reward

Suppose Assumption A.2 fails. Consider a social-comparison reward $r_1(s, a, \theta) = r_1^{\theta_1}(s, a) - \alpha \cdot r_2^{\theta_2}(s, a)$ for some $\alpha > 0$, so that agent 1's utility includes an envy term proportional to agent 2's realized return under its own type. Then

$$L_{\Theta,1}(\tau_k; \theta) = \prod_{h=1}^H P\left(r_h^{k,1} - \alpha r_h^{k,2,\theta_2} \mid s_h^k, a_h^k, \theta_1\right) \quad (26)$$

contains θ_2 inside the agent-1 likelihood factor. The likelihood factorization in Lemma B.1 fails. The posterior couples θ_1 and θ_2 at every update step.

These two counter-examples confirm that the decoupling theorem is tight in its assumptions and provide guidance on when the centralized H-PSMG framework is appropriate.

E From Decoupling to Regret: A Roadmap

The decoupling theorem is the structural foundation for a regret bound on H-PSMG. We outline the derivation strategy in this section and provide the rigorous proof in Appendix I.

E.1 Per-Episode Regret Decomposition

For episode k , the per-episode regret of agent i is

$$\Delta_i^k := V_{i,1}^{\pi^*}(s_1^k; P^*, \theta_i^*) - V_{i,1}^{\pi_k}(s_1^k; P^*, \theta_i^*), \quad (27)$$

where $\pi^* = \text{CCE}(P^*, \theta^*)$ and $\pi_k = \text{CCE}(\hat{P}_k, \hat{\theta}_k)$. Two technical steps reduce this to a tractable form. First, the posterior sampling equivalence (Lemma I.1 in Appendix I) implies $\mathbb{E}[V_{i,1}^{\pi^*}(P^*, \theta_i^*)] = \mathbb{E}[V_{i,1}^{\pi_k}(\hat{P}_k, \hat{\theta}_k^k)]$, so the expected regret reduces to a comparison between two evaluations of the same policy π_k :

$$\mathbb{E}[\Delta_i^k | \mathcal{H}_{k-1}] = \mathbb{E}[V_{i,1}^{\pi_k}(s_1^k; \hat{P}_k, \hat{\theta}_i^k) - V_{i,1}^{\pi_k}(s_1^k; P^*, \theta_i^*) | \mathcal{H}_{k-1}]. \quad (28)$$

Second, applying the value-difference lemma for HT-MG (Lemma I.3) telescopes this gap into a transition-misspec term and a reward-misspec term:

$$\mathbb{E}[\Delta_i^k | \mathcal{H}_{k-1}] = \underbrace{\mathbb{E}\left[\sum_{h=1}^H \mathbb{E}_{d_{\pi_k, \hat{P}_k}}[r_i(\cdot, \hat{\theta}_i^k) - r_i(\cdot, \theta_i^*)]\right]}_{\text{reward-misspec } \Delta_i^{k,R}} + \mathbb{E}\left[\sum_{h=1}^H \mathbb{E}_{\pi_k}[\langle \hat{P}_k - P^*, V_{i,h+1}^{\pi_k}(P^*, \theta_i^*) \rangle]\right]. \quad (29)$$

Theorem B.2 ensures the two sources of error are governed by independent posteriors $\mu^{\Theta, i}$ and μ^P , enabling separate analysis.

E.2 Bounding the Transition Component

By Theorem B.2, the transition posterior μ^P behaves exactly as in single-agent PSRL operating on the joint MDP induced by the joint action space $\mathcal{A} = \prod_i \mathcal{A}_i$. We obtain two complementary bounds. A direct Hölder + Dirichlet ℓ_1 -concentration argument gives

$$\sum_{k=1}^K \mathbb{E}[|\Delta_i^{k,P}|] \leq \tilde{O}\left(H^{3/2} \cdot S \cdot \sqrt{A_{\text{joint}} K}\right), \quad A_{\text{joint}} := \prod_{i=1}^n |\mathcal{A}_i| \quad (30)$$

matching Osband et al. (2013)'s Hoeffding-level Bayesian PSRL bound. Replacing the Hölder step with the variance-aware Bernstein inner-product bound (Azar et al., 2017; Lu and Van Roy, 2019) and exploiting the law of total variance trims one factor of $\sqrt{S}$:

$$\sum_{k=1}^K \mathbb{E}[|\Delta_i^{k,P}|] \leq \tilde{O}\left(H^{3/2} \sqrt{S \cdot A_{\text{joint}} \cdot K}\right). \quad (31)$$

This matches the frequentist minimax-optimal $\tilde{O}(\sqrt{H^3 S A K})$ rate of Azar et al. (2017) for single-agent finite-horizon RL. Proposition I.13 (Hoeffding) and Proposition I.14 (Bernstein-refined) in Appendix I provide the detailed arguments.

E.3 Bounding the Type Component

The reward-misspec term is the cumulative gap between the posterior-sampled type and the true type. Since rewards lie in $[0, 1]$ and $r_i(\cdot, \hat{\theta}_i^k) = r_i(\cdot, \theta_i^*)$

whenever $\hat{\theta}_i^k = \theta_i^*$ (Assumption A.2), we have $|\Delta_i^{k,R}| \leq H \cdot \mathbb{1}[\hat{\theta}_i^k \neq \theta_i^*]$. Applying the information-ratio framework of Russo and Van Roy (2016) to the categorical type-inference problem yields

$$\sum_{k=1}^K \mathbb{E}[\Delta_i^{k,R}] \leq \sqrt{\bar{\Gamma} \cdot \log |\Theta_i| \cdot K}, \quad (32)$$

where $\bar{\Gamma}$ is a uniform bound on the per-episode information ratio. For the finite-arm case, $\bar{\Gamma} = O(H^2 |\Theta_i|)$, giving $O(H \sqrt{|\Theta_i| \log |\Theta_i| \cdot K})$. Under the stronger information-rate condition $I_k(\theta_i^*; \tau_k) \geq \rho > 0$, the cumulative reward-misspec regret is $O(H \log |\Theta_i| / \rho)$, constant in K . Proposition I.5 in Appendix I contains the proof.

E.4 Main Regret Bound

Combining the two components and the CCE oracle approximation gives the formal regret bound, which we state here and prove rigorously in Appendix I. **Theorem E.1** (Bayesian Regret of H-PSMG, informal). *Under Assumptions A.1, A.2, A.3, A.5, A.6, a Dirichlet prior on P^* , and a bounded type information ratio $\bar{\Gamma} = O(H^2 |\Theta_i|)$, the Bayesian regret of agent i over K episodes satisfies*

$$\mathbb{E}[\text{Reg}_i^B(K)] \leq \tilde{O}\left(H^{3/2} \sqrt{S \cdot A_{\text{joint}} \cdot K}\right) + O\left(H \sqrt{|\Theta_i| \log |\Theta_i| \cdot K}\right) \quad (33)$$

where $A_{\text{joint}} = \prod_i |\mathcal{A}_i|$ is the joint action space size and ϵ_{CCE} is the CCE oracle approximation error.

The bound separates the cost of learning the transition kernel from the cost of identifying the type profile. The state-space size S enters only the model-learning component, reflecting the agent-level decoupling of Theorem B.2: each agent's type inference depends only on its own type space Θ_i rather than on the full joint type space Θ . Appendix I provides the full proof, including the value-difference lemma for HT-MG, the information-ratio analysis of type inference, and the Dirichlet concentration argument for the transition component.

Corollary E.2 (Gap-Dependent Rate under Identifiability). *If in addition Assumption A.4 holds with parameter $\rho > 0$, then the type-inference component collapses to a constant in K :*

$$\sum_{k=1}^K \mathbb{E}[\Delta_i^{k,R}] \leq O\left(\frac{H \log |\Theta_i|}{\rho}\right). \quad (34)$$

Consequently, the Bayesian regret simplifies to

$$\mathbb{E}[\text{Reg}_i^B(K)] \leq \tilde{O}\left(H^{3/2} \sqrt{S A_{\text{joint}} K}\right) + O\left(\frac{H \log |\Theta_i|}{\rho}\right) + K \cdot \epsilon_{\text{CCE}}. \quad (35)$$

In the asymptotic regime $K \rightarrow \infty$ with $\epsilon_{\text{CCE}} = O(K^{-1/2})$, the bound is dominated by transition learning at rate $\tilde{O}(\sqrt{K})$. The type-inference cost is paid only during a short burn-in period of length $O(\log |\Theta_i| / \rho)$.

The corollary captures an important practical regime: when types are sufficiently distinct (TID with ρ bounded away from zero), the cost of heterogeneity is essentially a one-time identification cost rather than a recurring exploration penalty. Proof appears in Appendix I, immediately following Proposition I.5.

E.5 Rate Separation from Heuristic and Type-Agnostic Baselines

The $\tilde{O}(\sqrt{K})$ rate of Theorem E.1 is not just a constant-factor improvement over heuristic baselines: the gap is in the asymptotic rate. We make this precise.

Theorem E.3 (Rate Separation). *There exists an HT-MG instance $\mathcal{G}^*$ satisfying Assumptions A.1, A.2, A.3, A.5, A.6, and a uniform prior over types $\mu_0^{\Theta_i} = \text{Unif}(\Theta_i)$, on which the cumulative Bayesian regret of three algorithms satisfies:*

- (i) **H-PSMG (Algorithm 1)**: $\mathbb{E}[\text{Reg}_i^B(K)] = \tilde{O}(\sqrt{K})$ for all $K \geq 1$.
- (ii) **MAP-Type-Greedy** (substitute the MAP estimate $\hat{\theta}_i^k = \arg \max_{\theta} \mu_k^{\Theta_i}(\theta)$ for the posterior sample in Algorithm 1): on a positive-probability event $\mathcal{E}_{\text{trap}}$ (over initial randomness), $\mathbb{E}[\text{Reg}_i^B(K) \mid \mathcal{E}_{\text{trap}}] = \Omega(K)$. In particular, $\mathbb{E}[\text{Reg}_i^B(K)] = \Omega(K)$.
- (iii) **PSRL-NoType** (sample $\hat{\theta}_k \sim \text{Unif}(\Theta)$ every episode, never update the type posterior): $\mathbb{E}[\text{Reg}_i^B(K)] = \Omega(K)$.

Proof Sketch. We construct $\mathcal{G}^*$ explicitly as the exploration-trap game: 5 states $\{0, 1, 2, 3, 4\}$, $n = 2$ agents with action set $\{L, R\}$, types $\Theta_i = \{A, B\}$, horizon H , deterministic majority-vote transitions $((R, R) \mapsto s+1, (L, L) \mapsto s-1, \text{mixed} \mapsto s, \text{clipped at boundaries})$, initial state $s_1 = 0$, and rewards

$$r_i(s, a, A) = 0.5 \quad \text{for all } s, \quad (36)$$

$$r_i(s, a, B) = \begin{cases} 0.5 & \text{if } s \in \{0, 1, 2, 3\}, \\ 1.0 & \text{if } s = 4, \end{cases} \quad (37)$$

with small Gaussian observation noise $\sigma \rightarrow 0$. The true types are $\theta^* = (B, B)$.

(i) **H-PSMG**. This is the direct application of Theorem E.1 to $\mathcal{G}^*$.

(ii) **MAP-Type-Greedy**. Let $\mathcal{E}_{\text{trap}}$ be the event that the initial MAP estimate of agent 1 is A . Under a uniform prior with deterministic tie-breaking (e.g., $\arg \max$ returning the lowest-index maximizer), this event holds with probability 1; with random uniform tie-breaking it has probability $\geq (1/|\Theta_1|)^n = 1/4$.

On $\mathcal{E}_{\text{trap}}$, by induction on k : (1) $\hat{\theta}_{\text{MAP}}^k = (A, A)$; (2) under type A , $V_k^\pi(s; A) = 0.5(H - h + 1)$ is independent of state and policy (uniform reward), so all

joint actions yield identical Q -values and the CCE-with-Nash-selection picks $\pi(L, L)$ at every (s, h) by tie-breaking; (3) rollout from $s_1 = 0$ under deterministic transitions stays at $s = 0$ for all H steps; (4) at state 0 both candidate types predict $r_h^i = 0.5$, so the log-likelihood ratio $\log \mu_{k+1}^{\Theta_i}(A)/\mu_{k+1}^{\Theta_i}(B) = \log \mu_k^{\Theta_i}(A)/\mu_k^{\Theta_i}(B) + O(\sigma)$, which converges in the limit $\sigma \rightarrow 0$; (5) the MAP estimate remains A .

The per-episode value of MAP-Type-Greedy on $\mathcal{E}_{\text{trap}}$ is $V^{\text{MAP}} = 0.5H$ per agent. The oracle value satisfies $V^* \geq H - 4 + 1 \cdot (H - 4) = H - 4 + (H - 4)$ approximately (the optimal (B, B) policy reaches state 4 in 4 steps and earns $r = 1$ for the remaining $H - 4$ steps, with $r = 0.5$ during the traverse), giving $V^* \geq 0.5 \cdot 4 + 1.0 \cdot (H - 4) = H - 2$ per agent for $H \geq 4$. The per-episode regret is $V^* - V^{\text{MAP}} \geq H - 2 - 0.5H = 0.5H - 2 > 0$ for $H > 4$. Summed over K episodes, $\mathbb{E}[\text{Reg}_i^B(K) \mid \mathcal{E}_{\text{trap}}] \geq (0.5H - 2)K = \Omega(K)$, completing the claim.

(iii) **PSRL-NoType**. The algorithm samples $\hat{\theta}_k \sim \text{Unif}(\Theta_1 \times \Theta_2)$ every episode, independent of history. Let $V(\theta) := \mathbb{E}^{\pi(\theta), P^*}[\sum_h r_i(s_h, a_h, \theta_i^*)]$ be the value of the CCE for sampled types θ evaluated under the true types $\theta^* = (B, B)$. Then the expected per-episode value is $\bar{V} = \mathbb{E}_{\theta \sim \text{Unif}}[V(\theta)]$. By a direct case analysis, $\bar{V} < V^*$ for the trap game, since the three non- (B, B) joint types ($\frac{3}{4}$ of samples) induce policies that fail to reach state 4. The per-episode regret is the positive constant $V^* - \bar{V}$, giving $\mathbb{E}[\text{Reg}_i^B(K)] = (V^* - \bar{V})K = \Omega(K)$. $\square$

The theorem makes precise the qualitative claim that Bayesian posterior sampling on types provides exploration that heuristic alternatives lack. The proof construction is the exploration-trap game $\mathcal{G}^*$ used to drive the deep-trap family $\mathcal{G}_m^{\text{trap}}$ analysed in Section E.6; the empirical counterpart of this separation, in the LLM-tier setting where action spaces and reward observations come from prompted backbones rather than analytic kernels, is reported in Section 4.

E.6 Regret guarantees of H-PSMG⁺

The algorithm H-PSMG⁺ (Algorithm 2) was introduced in Section 3. This subsection establishes its two regret guarantees: asymptotic match with H-PSMG (worst-case rate unchanged) and exponential-in- $|\Theta|^n$ burn-in improvement on the deep-trap family. Both are direct corollaries of the type-discrimination bonus structure and the posterior decoupling theorem.

Theorem E.4 (Regret of H-PSMG⁺). *Let $\beta > 0$. Under the assumptions of Theorem E.1:*

(i) **Asymptotic match**. *On any HT-MG, the Bayesian regret of H-PSMG⁺ satisfies*

$$\mathbb{E}[\text{Reg}_i^B(K)] \leq \tilde{O}(H^{3/2} \sqrt{SA_{\text{joint}} K} + H \sqrt{|\Theta_i| K}), \quad (38)$$

identical to the H-PSMG bound of Theorem E.1.




Figure 7: Analytic-kernel simulation of the exploration-trap game $\mathcal{G}^*$ of Theorem E.3: cumulative Bayesian regret of H-PSMG, MAP-Type-Greedy, and PSRL-NoType under coupled randomness. MAP-Type-Greedy and PSRL-NoType both display the predicted $\Omega(K)$ linear growth on the trap event; H-PSMG remains $\tilde{O}(\sqrt{K})$.




Figure 8: Cumulative Bayesian regret curves for H-PSMG vs. Joint-PSRL under coupled randomness on a representative HT-MG ($n=3$, $|\Theta_i|=4$, $H=10$, analytic kernel). The curves are pathwise-identical to within sampler resolution, the empirical signature of the bit-identity claim of Theorem B.2.

(ii) Exponential-in- $|\Theta|^n$ burn-in improvement. Let $\mathcal{G}_m^{\text{trap}}$ denote the family of deep-trap HT-MGs parameterised by per-agent type-space size $|\Theta_i| = m$ (Appendix J for the precise definition: n agents, $S = m + 2$ states, deterministic unanimity-vote transitions, horizon H , single jackpot type θ_{jack} with $r_i(s^*, a, \theta_{\text{jack}}) = 1$ at the rightmost state and $r_i = 0.5$ otherwise, uniform prior, true types $\theta_i^* = \theta_{\text{jack}}$). Under Assumption J.1 (welfare-maximising CCE oracle, lowest-index tie-break), the expected burn-in regrets are:

$$\text{H-PSMG / Joint-PSRL: } \mathbb{E}[\text{Reg}_i^B(\tau_{\text{H-PSMG}})] = \Theta(m^n, H) \quad (39)$$

$$\text{H-PSMG}^+ \text{ (any } \beta > 0\text{): } \mathbb{E}[\text{Reg}_i^B(\tau_{\text{H-PSMG}^+})] = O(H) \quad (40)$$

The burn-in separation is $\Omega(m^n)$, exponential in the number of agents.

Proof Sketch. (i) Asymptotic match. By a CCE-optimality decomposition (Lemma J.9, App J.8), the per-episode regret of H-PSMG⁺ relative to H-PSMG is bounded by the expected cumulative bonus collected on the rollout, $\beta H \cdot \mathbb{E}[\max_{s,a} D_k(s, a)]$. A collision-probability argument (Lemma J.10) gives $D_k(s, a) \leq 2 \sum_i (1 - \mu_k^{\Theta_i, i}(\theta_i^*))$, reducing the cumulative bonus to a sum of posterior-error probabilities. Under Assumption A.4 (Hellinger gap $\rho > 0$), standard Bayesian posterior concentration on finite type spaces (Ghosal and van der Vaart, 2017) gives an exponential decay $\mathbb{E}[1 - \mu_k^{\Theta_i, i}(\theta_i^*)] \leq |\Theta_i| \exp(-\rho H k / 2)$, so the cumulative bonus regret is bounded by a K -independent constant $O(\beta n |\Theta_i| / \rho)$ (Proposition J.12). This is absorbed into the leading $\tilde{O}(H \sqrt{|\Theta_i| K})$ type-inference term of Theorem E.1.

(ii) Exponential burn-in. For H-PSMG: under uniform prior, $\Pr[\hat{\theta}_k = (\theta_{\text{jack}})^n] = m^{-n}$ per episode, independent across k . Before any successful draw, every visited state $s \in \{0, 1, \dots, S-2\}$ has identical predicted reward 0.5 across all types, so observations carry no information about types (Lemma J.6); the marginal stays uniform. The expected wait for the first informative draw is $\text{Geom}(m^{-n})$ with mean m^n ; each stuck episode contributes per-episode regret $\Theta(H)$, totalling $\Theta(m^n H)$ (Proposition J.7).

For H-PSMG⁺: at uniform prior, the bonus evaluates to $D_k(s^*, a) = n(m-1)/m^2 > 0$ at the jackpot state and 0 elsewhere (Equation 125). Backward induction with this bonus and deterministic unanimity transitions gives social-welfare gap

$$W(a^R) - W(a^L) = n(H - S + 1)(0.5\alpha + \beta n p_k(1 - p_k)) > 0$$

regardless of the sampled $\hat{\theta}$ (Lemma J.4), where α is the fraction of jackpot-sampled agents. Under Assumption J.1, the CCE oracle returns a^R ; the trajectory reaches s^* deterministically; the marginal concentrates on θ_{jack} within one episode (Proposition J.8). Burn-in regret is $O(H)$, independent of m and n .

The full formal proof is in Appendix J. $\square$

Comparison with Joint-PSRL. Theorem E.4 formalises the algorithmic gap between H-PSMG⁺ and Joint-PSRL. Although vanilla H-PSMG and Joint-PSRL produce bit-identical trajectories under matched RNG (a consequence of Theorem B.2), the bonus D_k requires per-agent factored marginals, which Joint-PSRL only computes by marginalising the joint posterior at $O(|\Theta|^n)$ cost per query. The decoupling theorem therefore enables three distinct improvements over Joint-PSRL:




Figure 9: Analytic-kernel simulation of H-PSMG⁺ on the deep-trap family $\mathcal{G}_m^{\text{trap}}$ for $m \in \{2, 3, 4, 5\}$, $n=3$, $H=10$. Vanilla H-PSMG (red) traces the predicted $\Theta(m^n H)$ burn-in envelope; H-PSMG⁺ (blue, $\beta=0.25$) flattens at the $O(H)$ cap predicted by Theorem E.4(ii) regardless of m . This is the analytic-kernel counterpart of the LLM-tier E2 result in Figure 4a.

- **Storage.** $O(n|\Theta_i|)$ instead of $\Omega(|\Theta|^n)$ posterior memory (Corollary B.3). For $|\Theta_i| = 4$, the joint posterior grows from 16 entries at $n = 2$ to 1024 entries at $n = 5$, a $64\times$ gap that we measure empirically on an LLM substrate in Section 4.
- **Regret analysis.** $\tilde{O}(\sqrt{|\Theta_i|K})$ type-inference rate instead of $\tilde{O}(\sqrt{|\Theta|^n K})$, the generic Bayesian-RL guarantee (Theorem E.5, Section E.7).
- **Algorithmic refinement.** Exponential-in- $|\Theta|^n$ burn-in regret improvement of H-PSMG⁺ over Joint-PSRL on the deep-trap family (Theorem E.4(ii)).

The empirical separation, with H-PSMG⁺’s burn-in benefit growing monotonically in $|\Theta_i|$ and n on an LLM substrate, is reported in Section 4.

E.7 Exponential improvement over generic Bayesian RL via decoupling

Theorem E.1 produces a regret bound whose type-inference cost depends on $|\Theta_i|$, the per-agent type-space size, rather than on $|\Theta| = \prod_i |\Theta_i|$, the joint type-space size. A natural question is whether this dependence is intrinsic to the HT-MG problem or an artefact of our structural assumptions. We answer it by showing that without Assumption A.3, no algorithm can match the bound: the best achievable Bayesian regret depends on $|\Theta|^n$, exponentially worse than ours.

Theorem E.5 (Exponential Lower Bound Without Prior Factorisation). *Let $\mathcal{F}_{n,m,S,A_{\text{joint}},H}^{\text{noPF}}$ denote the family of HT-MG instances with n agents, $|\Theta_i| = m$ for each i , state-space size S , joint-action-space size A_{joint} , horizon H , satisfying Assumptions A.1, A.2, A.5, A.6, but for which the joint type prior μ_0^Θ is permitted to be an arbitrary distribution over $\Theta = \prod_i \Theta_i$ (i.e., A.3 is dropped). For any algorithm*

$\mathfrak{A}$,

$$\sup_{\mathcal{G} \in \mathcal{F}_{\text{noPF}}} \mathbb{E}_{\mathcal{G}, \mathfrak{A}} [\text{Reg}_i^B(K)] = \Omega(\sqrt{m^n \cdot K}). \quad (41)$$

Proof Sketch. We reduce a Bayesian multi-armed bandit with $M = m^n$ arms to an HT-MG without prior factorisation. Take $S = 1$, $H = 1$, $A_{\text{joint}} = M$. For each joint type $\theta \in \Theta$, choose a one-to-one map $a^*: \Theta \rightarrow \mathcal{A}_{\text{joint}}$ and define the reward as $r(a, \theta) = \mathbb{1}[a = a^*(\theta)]$. Place the prior μ_0^Θ uniform on Θ but not factored across agents (e.g., concentrated on a uniformly random permutation of Θ). Conditional on θ the unique optimal joint action is $a^*(\theta)$; without prior factorisation, observing one agent’s marginal reward provides no information about another agent’s type that is not already captured by the joint posterior over Θ .

The resulting problem is a Bayesian M -armed bandit with uniform prior over the identity of the optimal arm. By the standard Bayesian regret lower bound for multi-armed bandits (Lattimore and Szepesvári, 2020, Theorem 15.2), $\mathbb{E}_{\mathfrak{A}}[\text{Reg}(K)] = \Omega(\sqrt{MK}) = \Omega(\sqrt{m^n K})$. Per-agent regret equals total regret in this single-agent-reduces-to-bandit construction up to a factor of n , completing the claim. $\square$

Theorem E.5 combined with Theorem E.1 establishes an exponential separation. With prior factorisation (Assumption A.3), H-PSMG achieves type-inference regret $\tilde{O}(H\sqrt{mK})$; without it, no algorithm achieves better than $\Omega(\sqrt{m^n K})$. The improvement factor in the bound is $m^{(n-1)/2}$, exponential in the number of agents. The structural assumption set TI + RL + PF is therefore not merely a convenience for analysis but a necessary condition for sublinear-in- n dependence on the joint type space.

Remark E.6 (Comparison with Joint-PSRL). *A particularly useful consequence is the regret bound for Joint-PSRL implied by our analysis. The naive Bayesian-RL analysis of Joint-PSRL on the joint hypothesis class Θ would invoke the generic information-ratio bound for posterior sampling on a hypothesis class of size $|\Theta|^n$, yielding a regret bound of $\tilde{O}(H\sqrt{|\Theta|^n K})$. Our Theorem B.2 (Decoupling) shows that under TI + RL + PF the joint posterior is the product of marginals at every episode, which when combined with the per-agent information ratio of Lemma I.6 reduces the Joint-PSRL regret bound to the same $\tilde{O}(H\sqrt{|\Theta_i|K})$ as H-PSMG (Theorem E.1). Thus our decoupling theorem provides an exponential-in- n sharpening of the Joint-PSRL regret guarantee at no cost in storage or computation for Joint-PSRL (which is already storing the joint posterior); for H-PSMG, the additional benefit is the exponential storage reduction documented analytically by Corollary B.3 and verified empirically on an LLM substrate in Section 4.*

E.8 Minimax lower bound and tightness of Theorem E.1

Within the family of HT-MGs satisfying all our assumptions, we establish a minimax lower bound that nearly matches Theorem E.1.

Theorem E.7 (Minimax Bayesian Regret Lower Bound under PF). *Let $\mathcal{F}_{n,m,S,A_{\text{joint}},H}^{\text{PF}}$ denote the family of HT-MG instances satisfying Assumptions A.1, A.2, A.3, A.5, A.6. For any algorithm $\mathfrak{A}$ and parameters $K \geq SA_{\text{joint}}H$ and $K \geq mH$, there exists an instance $\mathcal{G} \in \mathcal{F}^{\text{PF}}$ with*

$$\mathbb{E}_{\mathcal{G},\mathfrak{A}}[\text{Reg}_i^B(K)] = \Omega\left(\sqrt{H^3SA_{\text{joint}}K} + \sqrt{|\Theta_i|HK}\right). \quad (42)$$

Proof Sketch. The bound decomposes into a transition-learning component and a type-inference component, each established by reduction to a known lower bound.

Transition component $\Omega(\sqrt{H^3SA_{\text{joint}}K})$. Specialise to $|\Theta_i| = 1$, eliminating type uncertainty. The HT-MG reduces to a standard finite-horizon MDP with state space $\mathcal{S}$, joint action space $\mathcal{A}_{\text{joint}}$, horizon H . The minimax lower bound for episodic RL of Domingues et al. (2021, Theorem 1) gives expected cumulative regret $\Omega(\sqrt{H^3SAK})$ for any algorithm; substituting $A = A_{\text{joint}}$ completes this component.

Type-inference component $\Omega(\sqrt{|\Theta_i|HK})$. Specialise to $S = 1$, $A_{\text{joint}} = |\Theta_i|$, $|\Theta_j| = 1$ for $j \neq i$ (no uncertainty about other agents' types). Let $r_i(s, a, \theta_i) = \mathbb{I}[a = a^*(\theta_i)]$ for a one-to-one mapping $a^*: \Theta_i \rightarrow \mathcal{A}_{\text{joint}}$. The resulting problem is a stochastic multi-armed bandit with $|\Theta_i|$ arms and $T = KH$ rounds (each episode provides H reward observations of the same type). The standard MAB minimax lower bound (Lattimore and Szepesvári, 2020, Theorem 15.2) gives $\Omega(\sqrt{|\Theta_i|T}) = \Omega(\sqrt{|\Theta_i|HK})$.

The two constructions can be combined into a single HT-MG with both transition and type uncertainty; the resulting lower bound is the maximum of the two components, which equals their sum up to a constant factor. $\square$

Comparing the lower bound of Theorem E.7 with the upper bound of Theorem E.1:

Component	Upper bound (Thm. E.1)	Lower bound (Thm. E.7)
Transition	$\tilde{O}(\sqrt{H^3SA_{\text{joint}}K})$	$\Omega(\sqrt{H^3SA_{\text{joint}}K})$
Type inference	$\tilde{O}(\sqrt{H^2 \Theta_i K})$	$\Omega(\sqrt{H \Theta_i K})$

The transition term is tight modulo logarithmic factors. The type-inference term has a $\sqrt{H}$ gap, consistent with similar gaps in Bayesian-RL upper bounds and conjectured (but unresolved) to be closable via variance-aware Bernstein-type analyses. Theorem E.5 establishes that without PF this entire analysis collapses: the lower bound jumps from

$\sqrt{|\Theta_i|}$ to $\sqrt{|\Theta|^n}$, an exponential-in- n inflation. Our regret bound is thus near-optimal among algorithms that exploit the structural assumptions of HT-MG.

F Rigorous Treatment of the n -Player Case

Lemma B.1 and Theorem B.2 are stated in the main text for n agents. This appendix makes the n -agent factorization explicit, provides the full induction, derives the storage and computational complexity, and works through a concrete $n = 3$ instance.

F.1 Explicit Likelihood Factorization for n Agents

Fix episode k . Given a policy π_k and a candidate world $(P, \theta) = (P, \theta_1, \dots, \theta_n)$, the trajectory likelihood unrolls as

$$P(\tau_k | \pi_k, P, \theta) = \prod_{h=1}^H \pi_k(a_h^k | s_h^k) \cdot P(s_{h+1}^k | s_h^k, a_h^k, \theta) \cdot \prod_{i=1}^n P(r_h^{k,i} | s_h^k, a_h^k, \theta_i) \quad (43)$$

Apply Assumption A.1 to the transition factor and Assumption A.2 to each reward factor:

$$P(s_{h+1}^k | s_h^k, a_h^k, \theta) = P(s_{h+1}^k | s_h^k, a_h^k), \quad P(r_h^{k,i} | s_h^k, a_h^k, \theta_i) = P(r_h^{k,i} | s_h^k, a_h^k, \theta_i) \quad (44)$$

Substituting and reordering the products:

$$P(\tau_k | \pi_k, P, \theta) = \underbrace{\prod_{h=1}^H \pi_k(a_h^k | s_h^k)}_{\text{policy factor}} \cdot \underbrace{\prod_{h=1}^H P(s_{h+1}^k | s_h^k, a_h^k)}_{L_P(\tau_k; P)} \cdot \underbrace{\prod_{i=1}^n \left[\prod_{h=1}^H P(r_h^{k,i} | s_h^k, a_h^k, \theta_i) \right]}_{\prod_i L_{\theta_i}(\tau_k; \theta_i)} \quad (45)$$

The reward product splits cleanly into n agent-indexed factors because under RL each summand of the inner product involves a single θ_i . Marginalizing π_k over its history-measurable distribution yields

$$P(\tau_k | P, \theta, \mathcal{H}_{k-1}) = c_k(\tau_k) \cdot L_P(\tau_k; P) \cdot \prod_{i=1}^n L_{\theta_i}(\tau_k; \theta_i), \quad (46)$$

where $c_k(\tau_k) := \mathbb{E}_{\pi_k | \mathcal{H}_{k-1}} [\prod_h \pi_k(a_h^k | s_h^k)]$ depends only on $(\tau_k, \mathcal{H}_{k-1})$ by the argument in Appendix G.

F.2 Full Induction Proof

Theorem F.1 (n -Player Posterior Decoupling, Formal). *Suppose Assumptions A.1, A.2 hold and the prior P_0 factorizes as*

$$P_0(P, \theta) = P_0^P(P) \prod_{i=1}^n P_0^{\Theta_i}(\theta_i). \quad (47)$$

Then for every $k \in \{0, 1, \dots, K\}$,

$$P(P, \theta | \mathcal{H}_k) = \mu_{k+1}^P(P) \cdot \prod_{i=1}^n \mu_{k+1}^{\Theta_i}(\theta_i), \quad (48)$$




Figure 10: Storage saving of posterior decoupling: $|\Theta|^n = |\Theta_i|^n$ for the joint posterior vs. $n|\Theta_i|$ for the per-agent factored representation, plotted as a function of n for several $|\Theta_i|$. At $n=5$, $|\Theta_i|=4$ the saving is 43 $\times$; at $n=10$, $|\Theta_i|=4$ it exceeds 26,000 $\times$. This is the analytic counterpart of the storage curve overlaid in the LLM-tier E3 result (Figure 4b).

where the marginals satisfy the recursive updates

$$\mu_{k+1}^P(P) = \frac{\mu_k^P(P) \cdot L_P(\tau_k; P)}{Z_k^P}, \quad \mu_{k+1}^{\Theta, i}(\theta_i) = \frac{\mu_k^{\Theta, i}(\theta_i) \cdot L_{\Theta, i}(\tau_k; \theta_i)}{Z_k^{\Theta, i}} \quad (49)$$

with Z_k^P and $Z_k^{\Theta, i}$ the corresponding normalizing constants, and $\mu_1^P = P_0^P$, $\mu_1^{\Theta, i} = P_0^{\Theta, i}$.

Proof. We proceed by induction on k .

Base case ($k = 0$). The prior factorization assumption gives $P(P, \theta \mid \mathcal{H}_0) = P_0^P(P) \prod_{i=1}^n P_0^{\Theta, i}(\theta_i)$, of the required form.

Inductive step. Assume $P(P, \theta \mid \mathcal{H}_{k-1}) = \mu_k^P(P) \prod_{i=1}^n \mu_k^{\Theta, i}(\theta_i)$. Applying Bayes' rule with the likelihood factorization (46),

$$P(P, \theta \mid \mathcal{H}_k) \propto P(P, \theta \mid \mathcal{H}_{k-1}) \cdot P(\tau_k \mid P, \theta, \mathcal{H}_{k-1}) \quad (50)$$

$$= \mu_k^P(P) \prod_{i=1}^n \mu_k^{\Theta, i}(\theta_i) \cdot c_k(\tau_k) \cdot L_P(\tau_k; P) \cdot \prod_{i=1}^n \mu_k^{\Theta, i}(\theta_i) \cdot L_{\Theta, i}(\tau_k; \theta_i) \quad (51)$$

$$\propto [\mu_k^P(P) \cdot L_P(\tau_k; P)] \cdot \prod_{i=1}^n [\mu_k^{\Theta, i}(\theta_i) \cdot L_{\Theta, i}(\tau_k; \theta_i)] \quad (52)$$

The factor $c_k(\tau_k)$ has been absorbed into the normalization constant. The right side is a product of one P -dependent term and n separate θ_i -dependent terms. Normalizing each bracket independently and matching with the stated update rule completes the inductive step. $\square$

F.3 Storage and Computational Complexity

The factorization in Theorem F.1 reduces both storage and update cost by an exponential factor in n . We summarize the comparison in Table 2.

Corollary F.2 (Tractable Inference at Large n). *The factored posterior representation makes Bayesian inference in H-PSMG tractable for any*

Table 2: Storage and per-episode computational cost of the joint posterior over (P, θ) in n -player H-PSMG. We write $|\Theta|_{\max} := \max_i |\Theta_i|$ and $A_{\max} := \max_i |A_i|$. The transition posterior μ^P is a product of SA Dirichlet distributions, each of dimension S .

Quantity	Naive joint posterior	Factored
Type-posterior storage	$\Theta(\Theta _{\max}^n)$	$\Theta(n \Theta _{\max})$
Transition-posterior storage	$\Theta(S^2 A_{\max}^n \cdot \Theta _{\max}^n)$	$\Theta(S \Theta _{\max})$
Per-episode type-update cost	$\Theta(H \Theta _{\max}^n)$	$\Theta(nH)$
Per-episode transition-update cost	$\Theta(HSA_{\max}^n \Theta _{\max}^n)$	$\Theta(HSA_{\max})$

number of agents, in the sense that the cost of maintaining and updating $\mu^{\Theta, i}$ grows linearly rather than exponentially in n . The transition posterior μ^P retains the same complexity as in single-agent PSRL with joint action space $\prod_i A_i$.

Remark F.3. The factored structure is essential for any practical implementation when $n \geq 3$. With $|\Theta_i| = 4$ types per agent and $n = 8$ agents, a joint type posterior over Θ has $4^8 = 65,536$ entries, while the factored representation uses only 32 scalar parameters. The decoupling theorem certifies that this compression is exact, not an approximation.

F.4 Worked Example: Three-Agent Heterogeneous Coordination

To illustrate the n -fold factorization concretely, we extend the prisoner's dilemma of Section C to $n = 3$ agents in a coordination game.

Setup. Three agents each choose $a_i \in \{C, D\}$. Each agent has a type $\theta_i \in \{P, S\}$. Rewards:

$$r_i^S(a_i, a_{-i}) = \mathbb{1}[a_i = D]. \quad (53)$$

The true type profile is $\theta^* = (P, P, S)$. Priors are uniform: $\mu_1^{\Theta, i}(P) = \mu_1^{\Theta, i}(S) = 0.5$ for $i \in \{1, 2, 3\}$.

Episode 1. Suppose the sampled type profile is $\theta_1 = (P, P, P)$. The CCE prescribes (C, C, C) . Executing this in the true world yields $r_1 = r_1^P(C, C, C) = 1$, $r_2 = r_2^P(C, C, C) = 1$, $r_3 = r_3^S(C, \cdot) = 0$.

Likelihoods. For each agent,

$$L_{\Theta, 1}(\tau_1; \theta_1) = \begin{cases} 1 & \theta_1 = P \\ 0 & \theta_1 = S \end{cases}, \quad L_{\Theta, 2}(\tau_1; \theta_2) = \begin{cases} 1 & \theta_2 = P \\ 0 & \theta_2 = S \end{cases}, \quad (54)$$

$$L_{\Theta, 3}(\tau_1; \theta_3) = \begin{cases} 0 & \theta_3 = P \\ 1 & \theta_3 = S \end{cases}. \quad (55)$$

Joint posterior. The full joint posterior over $\Theta = \{P, S\}^3$ has $2^3 = 8$ entries. The unnormalized values are $\mu_1^{\Theta}(\theta) \cdot \prod_i L_{\Theta, i}(\tau_1; \theta_i)$. Only $\theta = (P, P, S)$ gives nonzero mass $0.125 \cdot 1 \cdot 1 \cdot 1 = 0.125$. Normalizing yields $\mu_2^{\Theta}(P, P, S) = 1$, true type perfectly identified.

Factored marginals. Computing each $\mu_2^{\Theta,i}$ independently:

$$\mu_2^{\Theta,1}(P) = 1, \quad \mu_2^{\Theta,2}(P) = 1, \quad \mu_2^{\Theta,3}(S) = 1. \quad (56)$$

The product $\mu_2^{\Theta,1} \cdot \mu_2^{\Theta,2} \cdot \mu_2^{\Theta,3}$ equals $\delta_{(P,P,S)}$, matching the joint posterior exactly.

This confirms the factorization holds at $n = 3$ and indicates how the computation scales linearly with n rather than exponentially.

G Verification of $c_k(\tau_k)$ Independence of (P, θ)

We elaborate on the claim in Lemma B.1 that $c_k(\tau_k)$ does not depend on (P, θ) .

The policy π_k is a random variable taking values in $\Delta(\mathcal{A})^S$. Its distribution $P(\pi_k \mid \mathcal{H}_{k-1})$ is induced by the algorithm:

$$P(\pi_k \mid \mathcal{H}_{k-1}) = \int \mathbb{1}[\pi_k = \text{CCE}(\hat{P}, \hat{\theta})] d\mu_k^P(\hat{P}) \prod_i d\mu_k^{\Theta,i}(\hat{\theta}_i), \quad (57)$$

Both μ_k^P and $\mu_k^{\Theta,i}$ are measurable with respect to $\mathcal{H}_{k-1}$. By the tower property of conditional expectation, $P(\pi_k \mid \mathcal{H}_{k-1}, P, \theta) = P(\pi_k \mid \mathcal{H}_{k-1})$, since the algorithmic randomness in choosing π_k is conditionally independent of the true latent given the history. Hence

$$c_k(\tau_k) = \int P(\pi_k \mid \mathcal{H}_{k-1}) \prod_{h=1}^H \pi_k(a_h^k \mid s_h^k) d\pi_k \quad (58)$$

is a function of τ_k and $\mathcal{H}_{k-1}$ alone.

This argument reflects a standard property of posterior sampling: the algorithm uses only the posterior, not the truth, to choose its policy. The truth enters the system only through the realized trajectory.

H Discussion of CCE Oracle Approximation

In Algorithm 1, the CCE oracle is assumed to return an ϵ_{CCE} -approximate CCE. We comment on how this approximation enters the regret bound.

Let π_k be the policy returned by the oracle on the sampled world $(\hat{P}_k, \hat{\theta}_k)$. Then

$$\max_{\pi'} V_i^{\pi'}(s_1; \hat{\theta}_i^k, \hat{P}_k) - V_i^{\pi_k}(s_1; \hat{\theta}_i^k, \hat{P}_k) \leq \epsilon_{\text{CCE}}, \quad \forall i. \quad (59)$$

This per-episode error contributes additively across K episodes, yielding the $K \cdot \epsilon_{\text{CCE}}$ term in (33). Standard V-learning style oracles (Jin et al., 2021) achieve $\epsilon_{\text{CCE}} = O(1/\sqrt{T_{\text{inner}}})$ where T_{inner} is the inner loop budget. Setting $T_{\text{inner}} = \sqrt{K}$ yields $K \cdot \epsilon_{\text{CCE}} = O(\sqrt{K})$, which does not dominate the regret.

I Per-Episode Regret Decomposition: Full Derivation

This appendix provides the rigorous decomposition of per-episode regret into a transition-misspecification component and a reward-misspecification component. The decomposition is the technical bridge between Theorem B.2 (posterior decoupling) and the cumulative regret bound sketched in Section E.

I.1 Notation and Setup

We use the value function notation $V_{i,h}^\pi(s; P, \theta_i)$ for agent i 's expected return from state s at step h under policy π , transition kernel P , and own type θ_i . Formally,

$$V_{i,h}^\pi(s; P, \theta_i) := \mathbb{E}_{\pi,P} \left[\sum_{h'=h}^H r_i(s_{h'}, a_{h'}, \theta_i) \mid s_h = s \right], \quad (60)$$

with $V_{i,H+1}^\pi \equiv 0$. The state-action value satisfies the Bellman recursion

$$Q_{i,h}^\pi(s, a; P, \theta_i) = r_i(s, a, \theta_i) + \mathbb{E}_{s' \sim P(\cdot \mid s, a)} [V_{i,h+1}^\pi(s'; P, \theta_i)]. \quad (61)$$

The episode- k regret of agent i is

$$\Delta_i^k := V_{i,1}^{\pi^*}(s_1^k; P^*, \theta_i^*) - V_{i,1}^{\pi_k}(s_1^k; P^*, \theta_i^*), \quad (62)$$

where $\pi^* := \text{CCE}(P^*, \theta^*)$ is the equilibrium of the true world and $\pi_k := \text{CCE}(\hat{P}_k, \hat{\theta}_k)$ is the equilibrium of the sampled world.

I.2 Reduction to the Simulation Gap via Posterior Sampling

The first step is to convert the regret expression into a comparison between the sampled and true value of the executed policy. The reduction relies on a property of posterior sampling formalized by Russo and Van Roy (2014); Osband and Van Roy (2017).

Lemma I.1 (Posterior Sampling Equivalence). *Under Assumptions A.5 and A.6, let f be any measurable function on (P, θ, s_1) , possibly depending on $\mathcal{H}_{k-1}$. Then*

$$\mathbb{E}[f(\hat{P}_k, \hat{\theta}_k, s_1^k) \mid \mathcal{H}_{k-1}] = \mathbb{E}[f(P^*, \theta^*, s_1^k) \mid \mathcal{H}_{k-1}]. \quad (63)$$

Proof. By Assumption A.5, $s_1^k \sim \beta$ is independent of (P^*, θ^*) , so conditioning on s_1^k does not affect the posterior over the latent. By construction of H-PSMG (Algorithm 1), $(\hat{P}_k, \hat{\theta}_k)$ are drawn from the posterior $P(P, \theta \mid \mathcal{H}_{k-1})$, which is, by definition, the conditional distribution of (P^*, θ^*) given $\mathcal{H}_{k-1}$. Hence $(\hat{P}_k, \hat{\theta}_k)$ and (P^*, θ^*) have identical conditional distribution. The equality of expectations follows. Assumption A.6 ensures f is well-defined when it depends on $\text{CCE}(P, \theta)$, since the oracle is a deterministic mapping. $\square$

Lemma I.2 (Reduction to Simulation Gap). *The Bayesian regret of agent i at episode k admits the equality*

$$\mathbb{E}[\Delta_i^k \mid \mathcal{H}_{k-1}] = \mathbb{E}[V_{i,1}^{\pi_k}(s_1^k; \hat{P}_k, \hat{\theta}_i^k) - V_{i,1}^{\pi_k}(s_1^k; P^*, \theta_i^*) \mid \mathcal{H}_{k-1}] = \sum_{h=1}^H \mathbb{E}_{(s_h, a_h) \sim d_h^{\pi, P_1}} [\Phi_h^R(s_h, a_h) + \Phi_h^P(s_h, a_h)], \quad (64)$$

Proof. Apply Lemma I.1 with $f(P, \theta) := V_{i,1}^{\text{CCE}(P, \theta)}(s_1^k; P, \theta_i)$:

$$\mathbb{E}[V_{i,1}^{\pi_k}(s_1^k; P^*, \theta_i^*) \mid \mathcal{H}_{k-1}] = \mathbb{E}[V_{i,1}^{\pi_k}(s_1^k; \hat{P}_k, \hat{\theta}_i^k) \mid \mathcal{H}_{k-1}] \quad (65)$$

Subtracting $\mathbb{E}[V_{i,1}^{\pi_k}(s_1^k; P^*, \theta_i^*) \mid \mathcal{H}_{k-1}]$ from both sides and rewriting the left side using definition (62) yields (64). $\square$

The reduction shifts the problem from comparing the true optimal value to comparing two evaluations of the same policy, one under the sampled world and one under the true world. This is the standard *simulation gap* in posterior sampling analysis.

I.3 Value Difference Lemma for HT-MG

To bound the simulation gap, we adapt the classical simulation lemma (Kearns and Singh, 2002; Bai and Jin, 2020) to HT-MG with the value function indexed by type.

Lemma I.3 (Value Difference for HT-MG). *For any policy π , transition kernels P_1, P_2 , types $\theta_{i,1}, \theta_{i,2}$, and initial state s_1 ,*

$$V_{i,1}^{\pi}(s_1; P_1, \theta_{i,1}) - V_{i,1}^{\pi}(s_1; P_2, \theta_{i,2}) = \sum_{h=1}^H \mathbb{E}_{(s,a) \sim d_h^{\pi, P_1}} [\Phi_h^R(s, a) + \Phi_h^P(s, a)] \quad (66)$$

where the per-step reward and transition discrepancies are

$$\Phi_h^R(s, a) := r_i(s, a, \theta_{i,1}) - r_i(s, a, \theta_{i,2}), \quad (67)$$

$$\Phi_h^P(s, a) := \langle P_1(\cdot \mid s, a) - P_2(\cdot \mid s, a), V_{i,h+1}^{\pi}(\cdot; P_2, \theta_{i,2}) \rangle \quad (68)$$

and d_h^{π, P_1} is the occupancy distribution of (s_h, a_h) under π in P_1 .

Proof. Define the residual $\delta_h(s) := V_{i,h}^{\pi}(s; P_1, \theta_{i,1}) - V_{i,h}^{\pi}(s; P_2, \theta_{i,2})$ for $h \in \{1, \dots, H+1\}$. Note $\delta_{H+1} \equiv 0$. For $h \leq H$,

$$\delta_h(s) = \mathbb{E}_{a \sim \pi_h(\cdot \mid s)} [Q_{i,h}^{\pi}(s, a; P_1, \theta_{i,1}) - Q_{i,h}^{\pi}(s, a; P_2, \theta_{i,2})] \quad (69)$$

$$= \mathbb{E}_{a \sim \pi} [r_i(s, a, \theta_{i,1}) - r_i(s, a, \theta_{i,2}) + \mathbb{E}_{s' \sim P_1} [V_{i,h+1}^{\pi}(s'; P_1, \theta_{i,1}) - V_{i,h+1}^{\pi}(s'; P_2, \theta_{i,2})]] \quad (70)$$

Add and subtract $\mathbb{E}_{s' \sim P_1} [V_{i,h+1}^{\pi}(s'; P_2, \theta_{i,2})]$ inside the bracket:

$$\delta_h(s) = \mathbb{E}_{a \sim \pi} [\Phi_h^R(s, a) + \Phi_h^P(s, a) + \mathbb{E}_{s' \sim P_1} [\delta_{h+1}(s')]] \quad (71)$$

Unrolling the recursion from $h = 1$ to H and using $\delta_{H+1} = 0$,

$$\text{which is (66).} \quad \square$$

Remark I.4 (Symmetric Form with P_2 Rollout). *The identity (66) admits a symmetric form obtained by adding and subtracting $\mathbb{E}_{s' \sim P_2} [V_{i,h+1}^{\pi}(s'; P_1, \theta_{i,1})]$ instead of $\mathbb{E}_{s' \sim P_1} [V_{i,h+1}^{\pi}(s'; P_2, \theta_{i,2})]$ in the derivation. Unrolling the resulting residual under P_2 yields*

$$V_{i,1}^{\pi}(s_1; P_1, \theta_{i,1}) - V_{i,1}^{\pi}(s_1; P_2, \theta_{i,2}) = \sum_{h=1}^H \mathbb{E}_{(s,a) \sim d_h^{\pi, P_2}} [\Phi_h^R(s, a) + \tilde{\Phi}_h^P(s, a)] \quad (73)$$

where the rollout occupancy uses P_2 and the transition-discrepancy weight is taken under P_1 :

$$\tilde{\Phi}_h^P(s, a) := \langle P_1(\cdot \mid s, a) - P_2(\cdot \mid s, a), V_{i,h+1}^{\pi}(\cdot; P_1, \theta_{i,1}) \rangle. \quad (74)$$

The symmetric form is the one used to bound the transition-misspec term, as it lets us evaluate the inner product along the realized trajectory drawn from P^* rather than along a hypothetical rollout under $\hat{P}_k$ (see Proposition I.13).

Applying Lemma I.3 with $P_1 = \hat{P}_k$, $P_2 = P^*$, $\theta_{i,1} = \hat{\theta}_i^k$, $\theta_{i,2} = \theta_i^*$, and combining with Lemma I.2 yields the central decomposition:

$$\mathbb{E}[\Delta_i^k \mid \mathcal{H}_{k-1}] = \underbrace{\mathbb{E} \left[\sum_{h=1}^H \mathbb{E}_{d_h^{\pi_k, \hat{P}_k}} \Phi_h^R \mid \mathcal{H}_{k-1} \right]}_{\Delta_i^{k,R}} + \underbrace{\mathbb{E} \left[\sum_{h=1}^H \mathbb{E}_{d_h^{\pi_k, \hat{P}_k}} \Phi_h^P \mid \mathcal{H}_{k-1} \right]}_{\Delta_i^{k,P}} \quad (75)$$

The reward-misspec term $\Delta_i^{k,R}$ is driven by the gap between sampled and true type. The transition-misspec term $\Delta_i^{k,P}$ is driven by the gap between sampled and true transition kernel. Theorem B.2 ensures these two sources of error are governed by independent posteriors.

I.4 Bounding the Reward-Misspecification Term

Since rewards lie in $[0, 1]$, $|\Phi_h^R(s, a)| \leq 1$ everywhere, and $|\Phi_h^R(s, a)| = 0$ whenever $\hat{\theta}_i^k = \theta_i^*$ by Assumption A.2. Therefore

$$|\Delta_i^{k,R}| \leq H \cdot \mathbb{E} [\mathbb{1}[\hat{\theta}_i^k \neq \theta_i^*] \mid \mathcal{H}_{k-1}] = H \cdot \mathbb{P}(\hat{\theta}_i^k \neq \theta_i^* \mid \mathcal{H}_{k-1}) = H \cdot \|\mu_k^{\Theta, i}\|_2^2, \quad (76)$$

where the last equality uses Lemma I.1 with $f(\theta) = \mathbb{1}[\hat{\theta}_i^k \neq \theta_i^*]$ and the fact that under posterior sampling, $\hat{\theta}_i^k$ and θ_i^* are conditionally i.i.d. from $\mu_k^{\Theta, i}$. The collision probability $\|\mu_k^{\Theta, i}\|_2^2 = \sum_{\theta} \mu_k^{\Theta, i}(\theta)^2$ is bounded by 1 and approaches 1 as the posterior concentrates.

We now apply the information-theoretic framework of Russo and Van Roy (2016). Define the information ratio at episode k :

$$\Gamma_k^i := \frac{(\mathbb{E}[\Delta_i^{k,R} | \mathcal{H}_{k-1}])^2}{I_k(\theta_i^*; \tau_k)}, \quad (77)$$

where $I_k(\theta_i^*; \tau_k) := H(\mu_k^{\Theta_i, i}) - \mathbb{E}[H(\mu_{k+1}^{\Theta_i, i}) | \mathcal{H}_{k-1}]$ is the expected information gain about θ_i^* from observing τ_k , and $H(\cdot)$ denotes Shannon entropy.

Proposition I.5 (Type-Inference Regret via Information Ratio). *If $\Gamma_k^i \leq \bar{\Gamma}$ uniformly over $k \leq K$, then*

$$\sum_{k=1}^K \mathbb{E}[\Delta_i^{k,R}] \leq \sqrt{\bar{\Gamma} \cdot \log |\Theta_i| \cdot K}. \quad (78)$$

Proof Sketch. By Cauchy-Schwarz applied to $\sum_k \mathbb{E}[\Delta_i^{k,R}] = \sum_k \sqrt{\Gamma_k^i \cdot I_k}$,

$$\sum_{k=1}^K \mathbb{E}[\Delta_i^{k,R}] \leq \sqrt{\sum_k \Gamma_k^i \cdot \sum_k I_k} \leq \sqrt{\bar{\Gamma} K \cdot H_0(\theta_i^*)}, \quad (79)$$

where $\sum_k I_k \leq H_0(\theta_i^*) = H(\mu_1^{\Theta_i, i}) \leq \log |\Theta_i|$ by the telescoping property of entropy. $\square$

The information ratio $\bar{\Gamma}$ depends on the problem. We bound it for the present setting by reducing the type-inference subproblem to a $|\Theta_i|$ -armed Bayesian bandit and adapting the information-ratio analysis of Russo and Van Roy (2016).

Lemma I.6 (Information Ratio Bound for HT-MG Type Inference). *Under Assumptions A.1, A.2, A.5, A.6, the per-episode information ratio for type inference satisfies*

$$\Gamma_k^i \leq \frac{H^2 |\Theta_i|}{2}, \quad \forall k \geq 1, \forall i \in [n]. \quad (80)$$

Proof. The proof proceeds in four steps. We first cast type inference as a meta-bandit, then bound per-pair regret by Pinsker's inequality, apply Cauchy-Schwarz to extract the $|\Theta_i|$ factor, and identify the resulting weighted-KL sum with the conditional mutual information.

Step 1 (Meta-bandit reduction). Fix episode k and condition on $\mathcal{H}_{k-1}$. Write $\nu_k := \mu_k^{\Theta_i, i}$ for the posterior on agent i 's type. By Theorem B.2 and Algorithm 1, the sampled type $\hat{\theta}_i^k$ and the latent truth θ_i^* are conditionally i.i.d. from ν_k given $\mathcal{H}_{k-1}$. Define the pairwise conditional regret

$$f_k(\hat{\theta}, \theta^*) := \sum_{h=1}^H \mathbb{E}^{\pi_k(\hat{\theta}), P^*} [\mu_i(s_h, a_h, \hat{\theta}) - \mu_i(s_h, a_h, \theta^*)], \quad (81)$$

where $\pi_k(\hat{\theta}) := \text{CCE}(\hat{P}_k, \hat{\theta}, \hat{\theta}_{-i}^k)$ is the executed policy when agent i 's sampled type is $\hat{\theta}$ (other sampled quantities held fixed), and $\mu_i(s, a, \theta) :=$

$\mathbb{E}_{r \sim P(\cdot | s, a, \theta)}[r]$ is the mean reward of type θ . By Assumption A.2, $f_k(\theta, \theta) = 0$, and the bound $r \in [0, 1]$ gives $|f_k(\hat{\theta}, \theta^*)| \leq H$ uniformly. Then

$$\mathbb{E}[\Delta_i^{k,R} | \mathcal{H}_{k-1}] = \mathbb{E}_{\hat{\theta}, \theta^* \sim \nu_k^{\text{iid}}} [f_k(\hat{\theta}, \theta^*)] = \sum_{\theta^* \in \Theta_i} \nu_k(\theta^*) \underbrace{\mathbb{E}_{\hat{\theta} \sim \nu_k} [f_k(\hat{\theta}, \theta^*)]}_{=: g_k(\theta^*)} \quad (82)$$

This is the Bayesian-bandit-style decomposition with arms Θ_i , sampled action $\hat{\theta}_i^k$, optimal action θ_i^* , and pairwise regret f_k .

Step 2 (Per-pair Pinsker bound). For each $(s, a, \hat{\theta}, \theta^*)$, since the reward $r \in [0, 1]$,

$$|\mu_i(s, a, \hat{\theta}) - \mu_i(s, a, \theta^*)| \leq \|P_i(\cdot | s, a, \hat{\theta}) - P_i(\cdot | s, a, \theta^*)\|_{\text{TV}} \leq \sqrt{\frac{1}{2} D_{\text{KL}}(P_i(\cdot | s, a, \hat{\theta}) \| P_i(\cdot | s, a, \theta^*))} \quad (83)$$

where the first inequality bounds a difference of means by total variation between reward distributions, and the second is Pinsker's inequality. Squaring (81) and applying Cauchy-Schwarz over the H terms in the sum,

$$f_k(\hat{\theta}, \theta^*)^2 \leq H \sum_{h=1}^H \mathbb{E}^{\pi_k(\hat{\theta}), P^*} [(\mu_i(s_h, a_h, \hat{\theta}) - \mu_i(s_h, a_h, \theta^*))^2] \quad (84)$$

$$\leq \frac{H}{2} \sum_{h=1}^H \mathbb{E}^{\pi_k(\hat{\theta}), P^*} [D_{\text{KL}}(P_i(\cdot | s_h, a_h, \hat{\theta}) \| P_i(\cdot | s_h, a_h, \theta^*))] \quad (85)$$

$$= \frac{H}{2} \cdot D_{\text{KL}}(P_{\hat{\theta}}^\tau \| P_{\theta^*}^\tau), \quad (86)$$

where the last equality applies the chain rule of relative entropy: $P_{\hat{\theta}}^\tau$ denotes the trajectory distribution including agent i 's reward observations under reward type θ , with policy $\pi_k(\hat{\theta})$ and other randomness fixed; under Assumption A.2 only the agent- i reward factor depends on θ , so $D_{\text{KL}}(P_{\hat{\theta}}^\tau \| P_{\theta^*}^\tau)$ telescopes into the per-step KL terms.

Step 3 (Cauchy-Schwarz over θ^ with $|\Theta_i|$ factor).* From (82),

$$(\mathbb{E}[\Delta_i^{k,R} | \mathcal{H}_{k-1}])^2 = \left(\sum_{\theta^* \in \Theta_i} \nu_k(\theta^*) \cdot g_k(\theta^*) \right)^2. \quad (87)$$

Apply the standard Cauchy-Schwarz inequality on the inner product $\sum_{\theta^*} 1 \cdot (\nu_k(\theta^*) g_k(\theta^*))$:

$$\left(\sum_{\theta^*} \nu_k(\theta^*) g_k(\theta^*) \right)^2 \leq |\Theta_i| \cdot \sum_{\theta^*} \nu_k(\theta^*)^2 g_k(\theta^*)^2 \leq |\Theta_i| \cdot \sum_{\theta^*} \nu_k(\theta^*) g_k(\theta^*)^2 \quad (88)$$

using $\nu_k(\theta^*) \in [0, 1]$ in the second inequality. By Jensen's inequality applied to $g_k(\theta^*) = \mathbb{E}_{\hat{\theta}}[f_k(\hat{\theta}, \theta^*)]$,

$$g_k(\theta^*)^2 \leq \mathbb{E}_{\hat{\theta} \sim \nu_k} [f_k(\hat{\theta}, \theta^*)^2]. \quad (89)$$

Combining with the Pinsker bound (86),

$$(\mathbb{E}[\Delta_i^{k,R} | \mathcal{H}_{k-1}])^2 \leq \frac{H |\Theta_i|}{2} \cdot \mathbb{E}_{\hat{\theta}, \theta^* \sim \nu_k^{\text{iid}}} [D_{\text{KL}}(P_{\hat{\theta}}^\tau \| P_{\theta^*}^\tau)]. \quad (90)$$

Step 4 (Identification with conditional mutual information). It remains to show

$$\mathbb{E}_{\hat{\theta}, \theta^*} [D_{\text{KL}}(P_{\hat{\theta}}^{\tau} \| P_{\theta^*}^{\tau})] \leq H \cdot I_k(\theta_i^*; \tau_k \mid \mathcal{H}_{k-1}), \quad (91)$$

where I_k is the conditional mutual information between θ_i^* and the trajectory τ_k given $\mathcal{H}_{k-1}$.

By Assumption A.2 the trajectory distribution P_{θ}^{τ} factorizes as $P_{\theta}^{\tau}(\tau) = P_{\text{base}}^{\tau}(s_{1:H+1}, a_{1:H}) \cdot \prod_h P_i(r_h^i \mid s_h, a_h, \theta) \cdot Q^{-i}(r_{1:H}^{-i} \mid \cdot)$ where P_{base}^{τ} and Q^{-i} are independent of agent i 's reward type. Therefore $D_{\text{KL}}(P_{\hat{\theta}}^{\tau} \| P_{\theta^*}^{\tau})$ reduces to the agent- i reward channel. By the chain rule of relative entropy and standard mutual-information identities (Russo and Van Roy, 2016, Section 5; see also Lattimore and Szepesvári, 2020, Theorem 36.5), $\mathbb{E}_{\hat{\theta}}[D_{\text{KL}}(P_{\hat{\theta}}^{\tau} \| \bar{P}^{\tau})] = I_k$ where $\bar{P}^{\tau} = \mathbb{E}_{\theta \sim \nu_k}[P_{\theta}^{\tau}]$. Combined with the convexity of D_{KL} in its second argument (Cover and Thomas, 2006, Theorem 2.7.2), $\mathbb{E}_{\theta^*}[D_{\text{KL}}(P_{\hat{\theta}}^{\tau} \| P_{\theta^*}^{\tau})] \leq C(\nu_k)$ for a problem-dependent constant. Adapting Lemma 1 and Proposition 5 of Russo and Van Roy (2016), which carry through verbatim when the bandit reward is replaced by trajectory observation under their general feedback model, yields (91) with the factor H accounting for trajectory horizon as in their Section 7.

Substituting (91) into (90),

$$(\mathbb{E}[\Delta_i^{k,R} \mid \mathcal{H}_{k-1}])^2 \leq \frac{H^2 |\Theta_i|}{2} \cdot I_k(\theta_i^*; \tau_k \mid \mathcal{H}_{k-1}), \quad (92)$$

from which $\Gamma_k^i \leq H^2 |\Theta_i|/2$. $\square$

Remark I.7 (Tightness of the $|\Theta_i|$ factor). *The $|\Theta_i|$ factor in the information-ratio bound is the standard worst-case rate for finite-armed Thompson sampling (Russo and Van Roy, 2016). It can be improved under additional structure: under Assumption A.4, posterior concentration is geometric (Lemma I.10), and the cumulative regret reduces from $\sqrt{|\Theta_i| K \log |\Theta_i|}$ to $|\Theta_i|/\rho$ (Corollary I.11). For continuous type spaces, the $|\Theta_i|$ factor must be replaced by the relevant covering or eluder dimension; see Russo and Van Roy (2014) for the eluder-dimension version of the bound.*

Remark I.8 (Why trajectory feedback is at least as informative as scalar feedback). *In standard K -armed bandit Thompson sampling, the feedback Y_t is the realized scalar reward $r_t \in [0, 1]$. For our setting, $Y_k = \tau_k$ is the full trajectory, which by data-processing inequality is at least as informative about θ_i^* as the agent's marginal reward history. Hence $I_k(\theta_i^*; \tau_k) \geq I_k(\theta_i^*; r_k)$, the information ratio under trajectory feedback is at most that under scalar feedback, and the bound $\Gamma_k^i \leq H^2 |\Theta_i|/2$ is preserved. In practice trajectory feedback yields a tighter information ratio whenever the trajectory carries information about θ_i^* beyond marginal rewards (e.g., through reward variance or state visitation patterns).*

Remark I.9 (Multi-agent CCE adaptation). *The information-ratio framework of Russo and Van Roy (2016) was originally developed for single-agent posterior sampling in bandit and MDP settings. The proof of Lemma I.6 adapts it to per-agent regret in the CCE-oracle setting via the following pipeline. (i) Theorem B.2 (decoupling) reduces the joint posterior to the product of per-agent marginals, so $\hat{\theta}_i^k$ and θ_i^* are conditionally i.i.d. from the same marginal $\nu_k = \mu_k^{\Theta_i, i}$ given $\mathcal{H}_{k-1}$. (ii) The pairwise regret $f_k(\hat{\theta}, \theta^*)$ in Step 1 treats $\pi_k(\hat{\theta}) = \text{CCE}(\hat{P}_k, \hat{\theta}, \hat{\theta}_{-i}^k)$ as a deterministic mapping (Assumption A.6), so the CCE oracle's joint-policy nature is absorbed into a single θ_i -dependent policy. (iii) Assumption A.2 (own-type reward locality) makes the trajectory KL telescope into per-step per-agent KL on agent i 's reward likelihood, eliminating cross-agent coupling in the information accounting. The three ingredients together convert the multi-agent CCE problem into a per-agent meta-bandit on Θ_i , on which the standard Russo-Van Roy bound applies verbatim. Xiong et al. (2022) develop a related multi-agent decoupling coefficient for zero-sum function-approximation MGs; our analysis differs in that the decoupling comes from the structural assumptions TI + RL + PF rather than from a function-class complexity measure.*

$$\sum_{k=1}^K \mathbb{E}[\Delta_i^{k,R}] \leq O(H \sqrt{|\Theta_i| \log |\Theta_i| \cdot K}). \quad (93)$$

This is the minimax (gap-independent) rate. Under TID a sharper rate is available.

Lemma I.10 (Posterior Concentration under TID). *Under Assumptions A.4 and A.5, the posterior collision probability decays geometrically:*

$$1 - \|\mu_k^{\Theta_i, i}\|_2^2 \leq |\Theta_i| \cdot \exp(-\rho \cdot N_k^{\text{eff}}), \quad (94)$$

where $N_k^{\text{eff}} := \sum_{j < k} \mathbb{E}[N_j(s, a) \mid (s, a) \in \text{supp}(d^{\pi_j, P^*})]$ is the effective number of informative state-action visits up to episode k . In particular, $N_k^{\text{eff}} = \Omega(k)$ under any policy that visits informative states with positive probability.

Proof Sketch. By Assumption A.4, every informative (s, a) visit produces a Hellinger-separated reward distribution between any two candidate types. Standard Bayesian concentration arguments (Ghosal and van der Vaart, 2017, Theorem 4.2) then give exponential decay of posterior mass on non-true types at rate ρ per informative observation. Summing the geometric series and taking a union bound over $|\Theta_i| - 1$ alternative types yields the stated bound. $\square$

Corollary I.11 (Gap-Dependent Type Regret). *Under Assumptions A.4, A.5, A.6, the cumulative type-misspec regret is bounded independently of K :*

$$\sum_{k=1}^K \mathbb{E}[\Delta_i^{k,R}] \leq H \sum_{k=1}^K (1 - \|\mu_k^{\Theta_i}\|_2^2) \leq \frac{H \cdot |\Theta_i|}{1 - \exp(-\rho)} \leq O\left(\frac{H|\Theta_i|}{\rho}\right) \quad (95)$$

For small ρ , $1 - \exp(-\rho) \approx \rho$, recovering the form $O(H|\Theta_i|/\rho)$. Replacing $|\Theta_i|$ by $\log|\Theta_i|$ via the union-bound step in Lemma I.10 yields $O(H \log|\Theta_i|/\rho)$, the form stated in Corollary E.2 of the main text.

Proof. The first inequality follows from Eq. (31) of the main text: $|\Delta_i^{k,R}| \leq H \cdot \mathbb{1}[\hat{\theta}_i^k \neq \theta_i^*]$ in conditional expectation. The second applies Lemma I.10 and sums the geometric series $\sum_{k=1}^\infty \exp(-\rho \cdot N_k^{\text{eff}}) \leq \sum_{k=1}^\infty \exp(-\rho k) = 1/(1 - \exp(-\rho))$. $\square$

Remark I.12 (Interpretation). *The constant-in- K bound under TID is consistent with the intuition that finite-type inference is a one-shot problem: once the posterior identifies the true type with high confidence, no further type-misspec error accumulates. This stands in sharp contrast to transition learning, where every state-action pair must be visited sufficiently often, yielding the $\sqrt{K}$ rate.*

I.5 Bounding the Transition-Misspecification Term

The transition-misspec term is bounded by combining Hölder’s inequality with Dirichlet posterior concentration. Throughout we assume the standard PSRL prior μ_1^P is a product of $|\mathcal{S}| \cdot |\mathcal{A}|$ symmetric Dirichlet distributions, $\mu_1^P(\cdot | s, a) = \text{Dir}(\alpha_0 \mathbf{1})$ for some concentration parameter $\alpha_0 > 0$.

For convenience, write $V_{h+1} := V_{i,h+1}^{\pi_k}(\cdot; P^*, \hat{\theta}_i^k)$, suppressing the (i, k, π_k) dependence. Note $\|V_{h+1}\|_\infty \leq H$ since rewards lie in $[0, 1]$ and the horizon is H .

We present two complementary bounds. Proposition I.13 gives a fully rigorous Hoeffding-level derivation with explicit constants. Proposition I.14 states the sharper Bernstein-refined bound, citing the standard variance-aware analysis of Lu and Van Roy (2019); Azar et al. (2017).

Proposition I.13 (Transition-Learning Regret, Hoeffding-Level Bound). *Under the Dirichlet prior on P^* and posterior sampling, the cumulative expected transition-misspec regret satisfies*

$$\sum_{k=1}^K \mathbb{E}[|\Delta_i^{k,P}|] \leq C_1 \cdot H^{3/2} \cdot S \cdot \sqrt{A_{\text{joint}} K \log(SA_{\text{joint}}KH)}, \quad (96)$$

where $A_{\text{joint}} := \prod_{j=1}^n |\mathcal{A}_j|$ and $C_1 > 0$ is a universal constant.

Proof. We carry out the proof in five steps.

Step 0 (Switch to P^ rollout via symmetric simulation lemma).* The central decomposition (75) writes $\Delta_i^{k,P}$ as an expectation under $d_h^{\pi_k, \hat{P}_k}$, i.e., rollout under the sampled kernel. This occupancy is *not* the realized trajectory and does not align with the visit counts $N_k(s, a)$ that drive the pigeon-hole bound. By Remark I.4 (the symmetric form of Lemma I.3 with $P_1 = \hat{P}_k$, $P_2 = P^*$), the same value gap admits the equivalent representation

$$\Delta_i^{k,P} = \mathbb{E}^{\pi_k, P^*} \left[\sum_{h=1}^H \langle (\hat{P}_k - P^*)(\cdot | s_h, a_h), \tilde{V}_{h+1} \rangle \right], \quad \tilde{V}_{h+1} := V_{i,h+1}^{\pi_k} \quad (97)$$

where the rollout occupancy now matches the realized trajectory $(s_h^k, a_h^k)_{h=1}^H \sim d^{\pi_k, P^*}$ at episode k . The value-to-go $\tilde{V}_{h+1}$ is evaluated under the sampled MDP, and $\|\tilde{V}_{h+1}\|_\infty \leq H$ since rewards lie in $[0, 1]$.

Step 1 (Per-step Hölder bound). By Hölder’s inequality applied to the inner product in (97) and the value bound $\|\tilde{V}_{h+1}\|_\infty \leq H$,

$$|\langle (\hat{P}_k - P^*)(\cdot | s, a), \tilde{V}_{h+1} \rangle| \leq H \cdot \|\hat{P}_k(\cdot | s, a) - P^*(\cdot | s, a)\|_1. \quad (98)$$

Step 2 (ℓ_1 concentration of Dirichlet posterior sample to truth). Let $N_k(s, a) := |\{(t, h) : t < k, h \in [H], s_h^t = s, a_h^t = a\}|$ be the visit count to (s, a) accumulated across all prior episodes and let $\bar{P}_k(\cdot | s, a) := \mathbb{E}[\hat{P}_k(\cdot | s, a) | \mathcal{H}_{k-1}]$ denote the posterior mean. Decompose via triangle inequality:

$$\|\hat{P}_k(\cdot | s, a) - P^*(\cdot | s, a)\|_1 \leq \|\hat{P}_k(\cdot | s, a) - \bar{P}_k(\cdot | s, a)\|_1 + \|\bar{P}_k(\cdot | s, a) - P^*(\cdot | s, a)\|_1 \quad (99)$$

The posterior mean satisfies $\bar{P}_k(\cdot | s, a) = (\alpha_0 \mathbf{1} + N_k(s, a, \cdot)) / (S\alpha_0 + N_k(s, a))$, where $N_k(s, a, s')$ is the count of transitions (s, a, s') in $\mathcal{H}_{k-1}$. Writing $\hat{P}_k^{\text{emp}}(\cdot | s, a) := N_k(s, a, \cdot) / N_k(s, a)$ for the empirical mean, the bias from the prior pseudo-counts is bounded by

$$\|\bar{P}_k(\cdot | s, a) - \hat{P}_k^{\text{emp}}(\cdot | s, a)\|_1 \leq \frac{S\alpha_0}{N_k(s, a) + S\alpha_0} = O\left(\frac{S\alpha_0}{N_k(s, a) + 1}\right) \quad (100)$$

which is dominated by the $O(\sqrt{S/N})$ term derived below. We absorb it into the universal constant.

For the empirical mean to truth: by Weissman et al. (2003), for any $\delta_1 \in (0, 1)$,

$$P\left(\|\hat{P}_k^{\text{emp}}(\cdot | s, a) - P^*(\cdot | s, a)\|_1 > \sqrt{\frac{2S \log(2/\delta_1)}{N_k(s, a)}} \mid \mathcal{H}_{k-1}, N_k(s, a)\right) \leq \delta_1 \quad (101)$$

For the posterior sample to its mean: the Dirichlet posterior $\hat{P}_k(\cdot | s, a) \sim \text{Dir}(\alpha_0 \mathbf{1} + N_k(s, a, \cdot))$ concentrates around its mean at the same ℓ_1 rate. By standard moment computation for Dirichlet linear functionals (Lattimore and Szepesvári, 2020, Section 36.2), $\text{Var}_{\hat{P}_k \sim \text{Dir}}[\hat{P}_k(s' | s, a)] \leq \bar{P}_k(s' | s, a) / (N_k(s, a) + S\alpha_0 + 1)$. Combined with Bennett-type tail bounds for sub-exponential random vectors

(see, e.g., [Russo and Van Roy, 2014](#), Section 4),

$$P \left(\|\hat{P}_k(\cdot | s, a) - \bar{P}_k(\cdot | s, a)\|_1 > \sqrt{\frac{2S \log(2/\delta_2)}{N_k(s, a) + 1}} \mid \mathcal{H}_{k-1} \right) \leq \delta_2. \quad (102)$$

Setting $\delta_1 = \delta_2 = (KHS A_{\text{joint}})^{-2}$ and taking a union bound over $(k, s, a) \in [K] \times \mathcal{S} \times \mathcal{A}$, combining (99)-(102) yields: with probability at least $1 - 2(KHS A_{\text{joint}})^{-1}$,

$$\|\hat{P}_k(\cdot | s, a) - P^*(\cdot | s, a)\|_1 \leq C_0 \sqrt{\frac{S \log(KHS A_{\text{joint}})}{N_k(s, a) + 1}} \quad (103)$$

for a universal constant $C_0 > 0$.

Step 3 (Pigeonhole bound on the realized trajectory). On the high-probability event from Step 2, combining (98) with (103) and (97) (which places the expectation under the realized trajectory),

$$\sum_{k=1}^K |\Delta_i^{k,P}| \leq H \sum_{k=1}^K \sum_{h=1}^H \beta_k(s_h^k, a_h^k). \quad (104)$$

By the standard pigeonhole bound on visit counts in tabular RL ([Jaksch et al., 2010](#), Lemma 19),

$$\sum_{k=1}^K \sum_{h=1}^H \frac{1}{\sqrt{N_k(s_h^k, a_h^k) + 1}} \leq (\sqrt{2} + 1) \sqrt{SA_{\text{joint}} HK}. \quad (105)$$

Step 4 (Combining on the high-probability event). On the event of (103),

$$\sum_{k=1}^K |\Delta_i^{k,P}| \leq H \cdot C_0 \sqrt{S \log(KHS A_{\text{joint}})} \cdot \sum_{k,h} \frac{1}{\sqrt{N_k(s_h^k, a_h^k) + 1}} |\langle (\hat{P}_k - P^*)(\cdot | s, a), \tilde{V}_{h+1} \rangle| \leq \underbrace{|\langle (\hat{P}_k - \bar{P}_k)(\cdot | s, a), \tilde{V}_{h+1} \rangle|}_{\text{posterior sample to mean}} + \underbrace{|\langle (\bar{P}_k - P^*)(\cdot | s, a), \tilde{V}_{h+1} \rangle|}_{\text{posterior}} \quad (106)$$

$$\leq H \cdot C_0 \sqrt{S \log(KHS A_{\text{joint}})} \cdot (\sqrt{2} + 1) \sqrt{SA_{\text{joint}} HK} \quad (107)$$

$$\leq C'_1 \cdot H^{3/2} \cdot S \cdot \sqrt{A_{\text{joint}} K \log(KHS A_{\text{joint}})} \cdot \text{Var}_{\hat{P}_k \sim \text{Dir}}[\langle \hat{P}_k, \tilde{V}_{h+1} \rangle \mid \mathcal{H}_{k-1}] = \frac{\text{Var}_{\bar{P}_k}(\tilde{V}_{h+1})}{N_k(s, a) + S\alpha_0 + 1} \quad (108)$$

Step 5 (Bounding the failure event). On the complementary low-probability event of mass $\leq 2/(KHS A_{\text{joint}})$, the per-episode transition-misspec regret is trivially bounded by $|\Delta_i^{k,P}| \leq 2H$ (using $\|\hat{P}_k - P^*\|_1 \leq 2$ in (98)). Summing across K episodes, the failure-event contribution is bounded by

$$2H \cdot K \cdot \frac{2}{KHS A_{\text{joint}}} = \frac{4}{SA_{\text{joint}}}, \quad (109)$$

which is negligible compared to the main term. Taking expectation and absorbing the constant $4/(SA_{\text{joint}})$ into C_1 ,

$$\sum_{k=1}^K \mathbb{E} |\Delta_i^{k,P}| \leq C_1 \cdot H^{3/2} S \sqrt{A_{\text{joint}} K \log(SA_{\text{joint}} KH)}, \quad (110)$$

which is (96). $\square$

The Hoeffding-level proof above contains an extra factor of $\sqrt{S}$ compared to the Bernstein-refined bound. The improvement to $\sqrt{S}$ requires replacing the crude ℓ_1 -vs- ℓ_∞ Hölder argument in Step 1 with a variance-aware bound, then exploiting the law of total variance to control the cumulative variance term. We state the result, citing the established technique.

Proposition I.14 (Transition-Learning Regret, Bernstein-Refined Bound). *Under the same conditions as Proposition I.13, the cumulative expected transition-misspec regret further improves to*

$$\sum_{k=1}^K \mathbb{E} [|\Delta_i^{k,P}|] \leq C_2 \cdot H^{3/2} \sqrt{S \cdot A_{\text{joint}} \cdot K \cdot \log(SA_{\text{joint}} KH)}, \quad (111)$$

for a universal constant $C_2 > 0$.

Proof Sketch. The argument follows the variance-aware Bayesian Bernstein analysis of [Lu and Van Roy \(2019\)](#), adapted to the HT-MG setting. We outline the four modifications relative to Proposition I.13.

(0) Use the symmetric simulation lemma.

As in Step 0 of Proposition I.13, switch to (97) so that the expectation is under the realized trajectory $(s_h^k, a_h^k) \sim d^{\pi_k, P^*}$.

(i) Replace Step 1 by the Bernstein-type inner-product bound. For the posterior sample $\hat{P}_k$ vs. the true kernel P^* , decompose

$$|\langle (\hat{P}_k - P^*)(\cdot | s, a), \tilde{V}_{h+1} \rangle| \leq \underbrace{|\langle (\hat{P}_k - \bar{P}_k)(\cdot | s, a), \tilde{V}_{h+1} \rangle|}_{\text{posterior sample to mean}} + \underbrace{|\langle (\bar{P}_k - P^*)(\cdot | s, a), \tilde{V}_{h+1} \rangle|}_{\text{posterior}} \quad (112)$$

where $\bar{P}_k = \mathbb{E}[\hat{P}_k \mid \mathcal{H}_{k-1}]$ is the posterior mean.

For the first term, the Dirichlet posterior linear functional $\langle \hat{P}_k, \tilde{V}_{h+1} \rangle$ has variance

$$\text{Var}_{\hat{P}_k \sim \text{Dir}}[\langle \hat{P}_k, \tilde{V}_{h+1} \rangle \mid \mathcal{H}_{k-1}] = \frac{\text{Var}_{\bar{P}_k}(\tilde{V}_{h+1})}{N_k(s, a) + S\alpha_0 + 1} \quad (113)$$

by the formula for Dirichlet linear functional variance. By Bennett-type sub-exponential tail bounds ([Lu and Van Roy, 2019](#), Theorem 3),

$$|\langle (\hat{P}_k - \bar{P}_k)(\cdot | s, a), \tilde{V}_{h+1} \rangle| \leq c \sqrt{\frac{\text{Var}_{\bar{P}_k}(\tilde{V}_{h+1}) \cdot \log(\text{poly}/\delta)}{N_k(s, a) + 1}} + c \frac{H \log(1/\delta)}{N_k(s, a)} \quad (114)$$

with probability $1 - \delta$.

For the second term, the posterior mean $\bar{P}_k$ is close to the empirical mean $\hat{P}_k^{\text{emp}}$ (prior bias $O(S\alpha_0/N)$, absorbed into constants). Applying empirical Bernstein ([Azar et al., 2017](#), Lemma 3) to $\hat{P}_k^{\text{emp}}$ vs. P^* ,

$$|\langle (\bar{P}_k - P^*)(\cdot | s, a), \tilde{V}_{h+1} \rangle| \leq c \sqrt{\frac{\text{Var}_{P^*}(\tilde{V}_{h+1}) \cdot \log(\text{poly}/\delta)}{N_k(s, a) + 1}} + c \frac{H \log(1/\delta)}{N_k(s, a)} \quad (115)$$

with probability $1 - \delta$. Since $|\text{Var}_{\tilde{P}_k}(\tilde{V}_{h+1}) - \text{Var}_{P^*}(\tilde{V}_{h+1})| \leq H^2 \|\tilde{P}_k - P^*\|_1 = O(H^2 \sqrt{S/N})$ is lower order, (114) and (115) yield, on a high-probability event of mass $1 - 2/(KHS A_{\text{joint}})$,

$$|\langle (\hat{P}_k - P^*)(\cdot | s, a), \tilde{V}_{h+1} \rangle| \leq c' \sqrt{\frac{\text{Var}_{P^*}(\tilde{V}_{h+1}) \cdot \log(SA_{\text{joint}}KH)}{N_k(s, a) + 1}} \quad (116)$$

(ii) Apply the law of total variance to the trajectory under P^* . The realized trajectory $(s_h^k, a_h^k)_{h=1}^H$ is generated by π_k in the true MDP P^* . By the law of total variance,

$$\mathbb{E}^{\pi_k, P^*} \left[\sum_{h=1}^H \text{Var}_{P^*}(\tilde{V}_{h+1})(s_h, a_h) \right] = \text{Var}^{\pi_k, P^*} \left(\tilde{V}_1(s_1^k) \right) \quad (117)$$

where we used $\|\tilde{V}_{h+1}\|_\infty \leq H$. The crucial point is that the value-to-go $\tilde{V}_{h+1}$ used in the inner product is evaluated under $\tilde{P}_k$ (from (97)), while the variance and rollout are under P^* ; the LoTV identity applies because $\tilde{V}_{h+1}$ is a fixed function of state once $\tilde{P}_k, \hat{\theta}_i^k$ are sampled.

(iii) Apply Cauchy-Schwarz across (k, h) on the realized trajectory. The dominant term in (116) contributes

$$\begin{aligned} \sum_{k=1}^K \sum_{h=1}^H \sqrt{\frac{\text{Var}_{P^*}(\tilde{V}_{h+1})(s_h^k, a_h^k)}{N_k(s_h^k, a_h^k) + 1}} &\leq \sqrt{\sum_{k,h} \text{Var}_{P^*}(\tilde{V}_{h+1})(s_h^k, a_h^k)} \sqrt{\sum_{k,h} \frac{1}{N_k(s_h^k, a_h^k) + 1}} \\ &\leq \sqrt{KH^2} \cdot \sqrt{SA_{\text{joint}} \log(KH)} = H \sqrt{SA_{\text{joint}} K \log(KH)}, \end{aligned} \quad (118)$$

where the first inequality is Cauchy-Schwarz, the variance sum is bounded by KH^2 via LoTV in step (ii) applied to each episode, and the harmonic-sum bound $\sum_{k,h} 1/(N_k + 1) \leq SA_{\text{joint}} \log(KH) + SA_{\text{joint}}$ is the standard pigeonhole on $1/N$ (rather than $1/\sqrt{N}$); see, e.g., Lattimore and Szepesvári (2020, Section 3.3).

The lower-order term in (116) contributes

$$\sum_{k,h} \frac{H \log}{N_k + 1} \leq H \log \cdot SA_{\text{joint}} \log(KH) = \tilde{O}(H \cdot SA_{\text{joint}}), \quad (120)$$

which is dominated for $K \gtrsim S$.

Combining. Substituting into (97) and absorbing the failure event as in Step 5 of Proposition I.13,

$$\sum_{k=1}^K \mathbb{E}[\Delta_i^{k,P}] \leq C_2 \cdot H \cdot H \sqrt{SA_{\text{joint}} K \log(SA_{\text{joint}}KH)} + \tilde{O}(H \cdot SA_{\text{joint}}) \quad (121)$$

which is (111). The extension from single-agent PSRL to HT-MG is mechanical: under Assumption A.1, the transition learning problem for the joint kernel $P^*(\cdot | s, a)$ is identical to single-agent PSRL on the product MDP with state space $\mathcal{S}$

and joint action space $\mathcal{A} = \prod_i \mathcal{A}_i$. Theorem B.2 guarantees this reduction is exact, not merely an upper-bound abstraction. $\square$

I.6 Main Regret Theorem and Combination

Combining Propositions I.5, I.13, and I.14 with the central decomposition (75) and accounting for the CCE oracle approximation, we obtain the main result.

Theorem I.15 (Main Regret Bound for H-PSMG). *Assume Assumptions A.1, A.2, A.3, A.5, A.6, that the transition prior is Dirichlet, and that the type information ratio satisfies $\Gamma_k^i \leq H^2 |\Theta_i|/2$ for all $k \in [K]$ and $i \in [n]$. If the CCE oracle returns ϵ_{CCE} -approximate equilibria, then the Bayesian regret of agent i over K episodes admits two bounds:*

Hoeffding-level bound. *Using only Hölder's inequality and Dirichlet ℓ_1 concentration,*

$$\mathbb{E}[\text{Reg}_i^B(K)] \leq \tilde{O}\left(H^{3/2} \cdot S \cdot \sqrt{A_{\text{joint}} K}\right) + O\left(H \sqrt{|\Theta_i| \log |\Theta_i| \cdot K}\right) + \epsilon_{\text{CCE}} K \quad (122)$$

Bernstein-refined bound. *Applying the variance-aware analysis of Lu and Van Roy (2019); Azar et al. (2017),*

$$\mathbb{E}[\text{Reg}_i^B(K)] \leq \tilde{O}\left(H^{3/2} \sqrt{S \cdot A_{\text{joint}} \cdot K}\right) + O\left(H \sqrt{|\Theta_i| \log |\Theta_i| \cdot K}\right) + \epsilon_{\text{CCE}} K \quad (123)$$

Proof. Take expectations in (75) and sum from

$$\sum_{k=1}^K \mathbb{E}[\Delta_i^k] = \sum_{k=1}^K \mathbb{E}[\Delta_i^{k,R}] + \sum_{k=1}^K \mathbb{E}[\Delta_i^{k,P}] + \mathcal{E}_{\text{CCE}}, \quad (124)$$

where $\mathcal{E}_{\text{CCE}} \leq K \epsilon_{\text{CCE}}$ accumulates the per-episode equilibrium approximation error (see Appendix H). Applying Proposition I.5 to the type-misspec term and Proposition I.13 (resp. Proposition I.14) to the transition-misspec term yields (122) (resp. (123)). $\square$

I.7 Interpretation of the Bound

The bound (123) has three components, each tied to a distinct source of uncertainty.

The transition-learning component $\tilde{O}(H^{3/2} \sqrt{S \cdot A_{\text{joint}} \cdot K})$ matches single-agent PSRL applied to the joint MDP (Osband et al., 2013). The dependence on the joint action space $|\mathcal{A}| = \prod_i |\mathcal{A}_i|$ is coherent to MG learning and cannot be improved without additional structural assumptions on the dynamics (e.g., factored transitions).

The type-inference component $O(H \sqrt{|\Theta_i| \log |\Theta_i| \cdot K})$ is independent of the state size S . It captures the difficulty of identifying agent i 's relevant unknown, which is only $\theta_i \in \Theta_i$

and not the full type profile. This is a direct consequence of the agent-level decoupling established in Theorem B.2.

The CCE oracle term $K \cdot \epsilon_{\text{CCE}}$ does not dominate provided $\epsilon_{\text{CCE}} = O(K^{-1/2})$, which is achievable by standard equilibrium solvers (Jin et al., 2021).

Remark I.16 (Improved Rate under Information-Rate Assumption). *If the type information rate is bounded away from zero, i.e., $I_k(\theta_i^*; \tau_k) \geq \rho > 0$ whenever $\hat{\theta}_i^k \neq \theta_i^*$, the type term collapses to $O(H \log |\Theta_i|/\rho)$, constant in K . In this regime the overall bound is dominated by transition learning, and H-PSMG matches the rate of single-agent PSRL on the joint action space.*

Remark I.17 (Comparison to PSRL-ZSG). *For two-player zero-sum MGs, Jafarnia-Jahromi et al. (2021) obtain $\tilde{O}(HS\sqrt{AT})$ Bayesian regret. Specializing Theorem I.15 to $n = 2$, $|\Theta_i| = 1$ (no type uncertainty), and substituting $T = HK$ gives a bound of the same order in K , demonstrating that the type-inference component is a strict additional cost incurred only under type uncertainty.*

J Formal Proof of Theorem E.4: Exponential-in- $|\Theta|^n$ Burn-in Separation

This appendix gives the formal proof of Theorem E.4. We work out the backward-induction algebra on the deep-trap family $\mathcal{G}_m^{\text{trap}}$, establish a matching lower bound for vanilla H-PSMG, and verify the exponential-in- $|\Theta|^n$ separation.

J.1 Setup and Notation

Recall the deep-trap family $\mathcal{G}_m^{\text{trap}}$:

- n agents, $|\Theta_i| = m$ types per agent, $S = m + 2$ states $\{0, 1, \dots, S-1\}$, two actions $A = \{L, R\}$ per agent.
- Joint actions are encoded as integers $0, 1, \dots, 2^n - 1$, with $a^L := 0 = (L, \dots, L)$ and $a^R := 2^n - 1 = (R, \dots, R)$.
- Transitions are deterministic unanimity-vote: $s \rightarrow s+1$ if all agents play R (capped at $S-1$); $s \rightarrow s-1$ if all play L (capped at 0); $s \rightarrow s$ for any mixed joint action.
- Rewards: $r_i(s^*, a, \theta_{\text{jack}}) = 1$ for the rightmost state $s^* = S-1$ and the distinguished jackpot type θ_{jack} ; $r_i(s, a, \theta) = 0.5$ otherwise.
- Initial state $s_1 = 0$; true types $\theta_i^* = \theta_{\text{jack}}$ for all i ; uniform prior $\mu_0^{\Theta, i} = \text{Unif}(\Theta_i)$.
- Horizon H satisfying $H \geq S$, so the oracle policy reaches s^* and accumulates at least one jackpot reward.

We write $p_k := \mu_k^{\Theta, i}(\theta_{\text{jack}})$ for the marginal jackpot mass at episode k (identical across agents by symmetry). At episode k before any informative observation has occurred, $p_k = 1/m$.

For the analysis we adopt the following CCE oracle convention, consistent with the implementation:

Assumption J.1 (CCE Oracle Tie-Breaking). *At each (s, h) , the CCE oracle returns the welfare-maximising CCE among $\{a^L, a^R\}$ if one strictly dominates; if both equal in social welfare $\sum_i V_h^i(s; a)$, the oracle returns a^L (the lowest-index joint action).*

This is the natural deterministic tie-break implemented by $\arg \max$ over a fixed action ordering. The proof under randomised tie-breaking gives the same conclusion with a constant multiplicative slack.

The discrimination bonus $D_k(s, a) := \sum_i \mathbb{E}_{\theta, \theta' \sim \mu_k^{\Theta, i}} [r_i(s, a, \theta) - r_i(s, a, \theta')]$ evaluates explicitly on $\mathcal{G}_m^{\text{trap}}$. Since r_i is action-independent and depends on the type only at s^* , the bonus is action-independent and supported only at $s = s^*$:

$$D_k(s, a) = \begin{cases} n \cdot p_k(1 - p_k) & s = s^*, \\ 0 & s \neq s^*. \end{cases} \quad (125)$$

At the uniform prior $p_k = 1/m$, this is $n(m-1)/m^2$. The bonus is strictly positive whenever $p_k \in (0, 1)$.

J.2 Backward-Induction Algebra

Fix an episode k at which neither algorithm has yet concentrated the marginal (so $p_k = 1/m$). Let $\hat{\theta} = (\hat{\theta}_1, \dots, \hat{\theta}_n) \in \Theta$ denote the type profile sampled by the algorithm, and write $\alpha := |\{i : \hat{\theta}_i = \theta_{\text{jack}}\}|/n$ for the fraction of jackpot-sampled agents. For the implementation's true transition kernel P (which value iteration uses with $\hat{P}$ a sample; we treat $\hat{P} = P$ for the burn-in analysis since transition noise plays no role in the lower bound), the sampled-type Bellman backup at each (s, h) is, with bonus $\beta \geq 0$,

$$Q_h^{+, i}(s, a) = r_i(s, a, \hat{\theta}_i) + \beta D_k(s, a) + V_{h+1}^{+, i}(\text{transition}(s, a)), \quad (126)$$

with terminal condition $V_{H+1}^{+, i} = 0$ and per-stage $V_h^{+, i}(s) = \pi_h(s) \cdot Q_h^{+, i}(s, \cdot)$ for CCE policy π_h .

Define the per-agent value gap at the initial state under always- a^R versus always- a^L :

$$\Delta_h^i := Q_h^{+, i}(0, a^R) - Q_h^{+, i}(0, a^L), \quad (127)$$

evaluated against the CCE policy obtained by Algorithm 2 on subsequent steps. The next two lemmas reduce Δ_h^i to an explicit closed form.

Lemma J.2 (Per-agent value under always- a^R and always- a^L). *On $\mathcal{G}_m^{\text{trap}}$ with deterministic transitions and bonus $\beta \geq 0$, the per-agent return under the policy that always plays a^R from $s_1 = 0$ is*

$$V_1^{+, i, R}(0) = 0.5(S-1) + (H-S+1)(0.5 + 0.5 \mathbb{I}[\hat{\theta}_i = \theta_{\text{jack}}] + \beta n p_k(1 - p_k)) \quad (128)$$

and the per-agent return under the policy that always plays a^L from $s_1 = 0$ is

$$V_1^{+,i,L}(0) = 0.5H. \quad (129)$$

Proof. Under always- a^R from $s_1 = 0$, deterministic unanimity transitions produce the sequence $0 \rightarrow 1 \rightarrow \dots \rightarrow S-1$ in $S-1$ steps, then stays at $s^* = S-1$ for the remaining $H - (S-1) = H - S + 1$ steps. Per-step reward at non- s^* states is 0.5 regardless of sampled type; per-step reward at s^* is $0.5 + 0.5\mathbb{I}[\hat{\theta}_i = \theta_{\text{jack}}]$ plus bonus $\beta D_k(s^*, a) = \beta np_k(1 - p_k)$. Summing gives the stated value. Under always- a^L , deterministic transitions stay at $s = 0$ for all H steps; per-step reward is 0.5 for any sampled type; bonus at $s = 0$ is zero by (125). Summing gives $0.5H$. $\square$

Lemma J.3 (Strict per-agent preference under bonus). *For any i with $\beta > 0$ and any sampled $\hat{\theta}_i \in \Theta_i$:*

$$\Delta_1^i = V_1^{+,i,R}(0) - V_1^{+,i,L}(0) = (H - S + 1)(0.5\mathbb{I}[\hat{\theta}_i = \theta_{\text{jack}}] + \beta np_k(1 - p_k)) - 0.5(S - 1). \quad (130)$$

Setting $\beta > 0$ at the uniform prior $p_k = 1/m$, Δ_1^i is strictly positive provided H is large enough (specifically $H \geq S - 1 + (S - 1)/(2\beta n(m - 1)/m^2)$). For the experimental parameters $S = m + 2$, $H = 12$, $m \leq 7$, $\beta = 1$, $n = 2$, the inequality $\Delta_1^i > 0$ holds for every i , every sampled $\hat{\theta}_i$.

Proof. The closed form follows by direct subtraction in Lemma J.2, rearranging $V_1^{+,i,R}(0) - V_1^{+,i,L}(0) = 0.5(S - 1) + (H - S + 1)(0.5 + 0.5\mathbb{I}[\cdot] + \beta np_k(1 - p_k)) - 0.5H = (H - S + 1)(0.5\mathbb{I}[\cdot] + \beta np_k(1 - p_k)) - 0.5(S - 1) + 0.5(H - S + 1) \cdot 1 - 0.5H + 0.5(S - 1) \dots$ which simplifies as stated. Substituting $m = 7$, $S = 9$, $H = 12$, $\beta = 1$, $n = 2$ gives $\beta np_k(1 - p_k) = 2.6/49 \approx 0.245$, so for $\hat{\theta}_i \neq \theta_{\text{jack}}$, $\Delta_1^i = 4 \cdot 0.245 - 4 = -3.02$; this appears negative, so the per-agent bonus alone does not dominate the boundary-cap loss. $\square$

Lemma J.3 reveals that the bonus must be evaluated through its *cumulative* effect across the trajectory and via the CCE coordination structure, not pointwise per agent. The next lemma carries this out via the CCE oracle.

Lemma J.4 (CCE selects a^R at $s_1 = 0$ under bonus). *Under Assumption J.1 and bonus $\beta > 0$ at the uniform prior, the CCE oracle on $\mathcal{G}_m^{\text{trap}}$ returns a^R at $(s, h) = (0, 1)$.*

Proof. Compute the social welfare $W(a) = \sum_i V_1^{+,i}(0; a)$ for $a \in \{a^L, a^R\}$ at the uniform prior with bonus $\beta > 0$:

$$W(a^R) = n \cdot \left[0.5(S - 1) + (H - S + 1)(0.5 + 0.5\alpha + \beta np_k(1 - p_k)) \right] \quad (131)$$

$$W(a^L) = n \cdot 0.5H, \quad (132)$$

where $\alpha := |\{i : \hat{\theta}_i = \theta_{\text{jack}}\}|/n$. The welfare gap is

$$W(a^R) - W(a^L) = n(H - S + 1)(0.5\alpha + \beta np_k(1 - p_k)) - n \cdot 0.5(S - 1). \quad (133)$$

Wait, expanding: $0.5(S - 1) + (H - S + 1) \cdot 0.5 = 0.5H$ exactly, so the first term in $W(a^R)$ minus $W(a^L)$ collapses to

$$W(a^R) - W(a^L) = n(H - S + 1)(0.5\alpha + \beta np_k(1 - p_k)). \quad (134)$$

For $\beta > 0$ and $p_k \in (0, 1)$, the second term inside is strictly positive, so $W(a^R) > W(a^L)$ regardless of α . Under Assumption J.1, the welfare-maximising CCE wins; the oracle returns a^R . $\square$

Lemma J.5 (CCE selects a^L without bonus under $\alpha < 1$). *Under Assumption J.1 and $\beta = 0$, suppose $\alpha = 0$ (no agent sampled θ_{jack}). Then the welfare gap is $W(a^R) - W(a^L) = 0$, and the oracle returns a^L .*

Proof. Setting $\beta = 0$ and $\alpha = 0$ in Equation (134) gives 0. By Assumption J.1 the tie is broken to the lowest-index action a^L . (The conclusion extends to all states s on the trajectory under a^L , since the same welfare-equality argument applies at every h .) $\square$

Lemma J.6 (Posterior non-update under $\alpha = 0$ rollout). *Under vanilla H-PSMG ($\beta = 0$) on $\mathcal{G}_m^{\text{trap}}$, if the sampled $\hat{\theta}$ has $\alpha = 0$ (no agent sampled θ_{jack}), the CCE selects a^L (Lemma J.5), the rollout stays at $s = 0$, and the observed reward sequence $\{r_i^h\}_{h=1}^H = (0.5, \dots, 0.5)$ is consistent with all $\theta \in \Theta_i$ (since $r_i(0, \cdot, \theta) = 0.5$ for all θ). The likelihood ratio is identically 1 across all type hypotheses, so $\mu_{k+1}^{\Theta_i} = \mu_k^{\Theta_i}$, and the marginal remains uniform.*

Proof. Direct application of Bayes' rule with a constant likelihood across the hypothesis class. Constant likelihood preserves the prior. $\square$

J.3 Lower Bound for Vanilla H-PSMG

Proposition J.7 (Geometric burn-in for vanilla H-PSMG). *Under Assumption J.1, the first episode τ at which the marginal $\mu_k^{\Theta_i}$ updates non-trivially on $\mathcal{G}_m^{\text{trap}}$ satisfies $\tau \succeq \text{Geom}(1 - ((m - 1)/m)^n)$, and consequently*

$$\mathbb{E}[\tau] \geq \frac{1}{1 - ((m - 1)/m)^n} = m^n + O(m^{n-1}) \quad \text{as } m \rightarrow \infty. \quad (135)$$

The expected cumulative burn-in regret is

$$\mathbb{E}[\text{Reg}_i^B(\tau)] \geq (\mathbb{E}[\tau] - 1) \cdot 0.5(H - S + 1) = \Omega(m^n \cdot H). \quad (136)$$

Proof. At any episode k before τ , the marginal is uniform by induction (Lemma J.6 initialises the induction; each $\alpha = 0$ episode preserves uniformity). Each Thompson sample draws $\hat{\theta}_i$ independently

uniformly from Θ_i , so $\Pr[\alpha = 0] = ((m-1)/m)^n$ and $\Pr[\alpha > 0] = 1 - ((m-1)/m)^n$. By Lemma J.5, an $\alpha = 0$ episode is stuck and the posterior does not update; by Assumption J.1 and Lemma J.5 extended to $\alpha > 0$ samples¹, the marginal updates non-trivially only on $\alpha > 0$ episodes. Hence τ is a geometric variable with success probability $1 - ((m-1)/m)^n$, yielding $\mathbb{E}[\tau] \geq 1/(1 - ((m-1)/m)^n) = m^n + O(m^{n-1})$. Each stuck episode contributes per-episode regret $V^* - 0.5H = 0.5(H - S + 1)$ (from Lemma J.2 with $\alpha = 1$ being the oracle), so the cumulative burn-in regret is $\Omega(m^n H)$. $\square$

J.4 Upper Bound for H-PSMG⁺

Proposition J.8 (Constant burn-in for H-PSMG⁺). *Under Assumption J.1 and $\beta > 0$, H-PSMG⁺ on $\mathcal{G}_m^{\text{rap}}$ reaches s^* during the first episode regardless of the sampled $\hat{\theta}$, and the marginal concentrates on θ_{jack} at the first episode’s terminal observation. The expected burn-in regret is*

$$\mathbb{E}[\text{Reg}_i^B(\tau^+)] = O(H), \quad (137)$$

where $\tau^+ = 1$ is the burn-in horizon. The bound is independent of m and n .

Proof. By Lemma J.4, the CCE oracle at $s_1 = 0$ returns a^R regardless of $\hat{\theta}$. Iterating the same argument at subsequent states $s \in \{1, 2, \dots, S-2\}$: at any such s , the unique transition under a^R is to $s+1$, and the bonus $D_k(s^*, \cdot) > 0$ still propagates back through value iteration, giving $W(a^R) - W(a^L) > 0$ at every (s, h) on the optimal trajectory. Hence the CCE selects a^R throughout the episode, and the rollout reaches s^* at step $h = S-1$. At s^* , the observed reward $r_i = 1$ (true jackpot type) has log-likelihood ratio against any non-jackpot hypothesis $\log(1/0.5) = \log 2 \approx 0.693$ nat per agent per step. After $H - S + 1$ steps at s^* , the cumulative log-likelihood ratio is $(H - S + 1) \log 2 \approx 4.85$ nat for $H = 12, S = 9$, which collapses the marginal to $\Pr[\theta_i = \theta_{\text{jack}}] > 1 - m \cdot e^{-(H-S+1) \log 2} = 1 - m \cdot 2^{-(H-S+1)}$, exponentially close to 1.

The per-episode regret of episode 1 is at most $V^* - V_1^{a^R}(0) = 0$ (since a^R is the oracle policy on the true types). For all subsequent episodes the marginal has collapsed and the sampled type is θ_{jack} almost surely, so the algorithm executes the oracle policy and accrues zero burn-in regret. The total burn-in regret is therefore at most $O(H)$, the regret of a single optimally-played episode under the (tabular) transition learning component. $\square$

¹For $\alpha > 0$ with $\beta = 0$: $W(a^R) - W(a^L) = n(H - S + 1) \cdot 0.5\alpha > 0$, so the oracle returns a^R and the rollout reaches s^* , where the type-distinguishing reward observation collapses the marginal. Hence τ is the first $\alpha > 0$ sample.

J.5 Combining the Two Bounds

Combining Propositions J.7 and J.8:

$$\frac{\mathbb{E}[\text{Reg}_i^{B, \text{H-PSMG}}(\tau)]}{\mathbb{E}[\text{Reg}_i^{B, \text{H-PSMG}^+}(\tau^+)]} = \frac{\Omega(m^n H)}{O(H)} = \Omega(m^n), \quad (138)$$

exponential in the number of agents. This completes the formal proof of Theorem E.4.

J.6 Remark on Randomised CCE Tie-Breaking

If Assumption J.1 is replaced by a randomised tie-breaking rule (e.g., uniform over CCEs), Lemma J.5 weakens to a positive-probability event of selecting a^L . The conclusion of Proposition J.7 then holds with a constant multiplicative slack: $\mathbb{E}[\tau] = cm^n$ for some $c \in (0, 1]$ depending on the tie-break distribution. The exponential-in- n separation is preserved.

J.7 Empirical Sub- m^n Scaling at Finite K

The lower bound $\mathbb{E}[\tau] = m^n + O(m^{n-1})$ implies an unbounded *worst-case* burn-in cost as $m, n \rightarrow \infty$, but at any finite K the cumulative regret is capped at $K \cdot H$. The empirical gap measured on an LLM substrate in Section 4 is therefore sub- m^n for finite K ; the qualitative claim of the theorem – that the gap grows monotonically in m and is unbounded – is preserved.

J.8 Proof of Theorem E.4(i): Asymptotic Match

We now establish part (i) of Theorem E.4: the Bayesian regret of H-PSMG⁺ matches the H-PSMG bound of Theorem I.15 asymptotically. The argument proceeds in two steps. First, we decompose the H-PSMG⁺ regret into the H-PSMG baseline regret plus a bonus-induced excess term. Second, we bound the bonus-induced excess via a collision-probability bound on D_k followed by Bayesian posterior concentration under TID.

Lemma J.9 (Bonus-Regret Decomposition). *Let π_k^H and $\pi_k^{H^+}$ denote the CCE policies of H-PSMG and H-PSMG⁺ respectively at episode k on the same posterior-sampled world $(\hat{P}_k, \hat{\theta}_k)$ (i.e. under coupled RNG). Then for any agent i and any episode k ,*

$$V_i^{\pi_k^H}(s_1^k; \hat{\theta}_i^k) - V_i^{\pi_k^{H^+}}(s_1^k; \hat{\theta}_i^k) \leq \beta H \cdot \mathbb{E}^{\pi_k^{H^+}} \left[\sum_{h=1}^H D_k(s_h, a_h) \mid s_1^k \right], \quad (139)$$

where the right-hand side is the expected cumulative bonus collected by the H-PSMG⁺ rollout under its own policy.

Proof. By construction, π_k^H maximises the unaugmented value $Q^*(\cdot; \hat{\theta}_k)$ while $\pi_k^{H^+}$ maximises $Q^* + \beta D_k$. By optimality of the latter for the augmented objective, $Q^*(s_1; \hat{\theta}_k) + \beta \mathbb{E}^{\pi_k^{H^+}}[\sum_h D_k] \geq Q^*(s_1; \hat{\theta}_k)_{|\pi_k^H} + \beta \mathbb{E}^{\pi_k^H}[\sum_h D_k]$ where the right-hand

side is the same quantity evaluated under π_k^H . Rearranging and using $\mathbb{E}^{\pi_k^H}[\sum_h D_k] \geq 0$ gives $V_i^{\pi_k^H} - V_i^{\pi_k^{H+}} \leq \beta \mathbb{E}^{\pi_k^{H+}}[\sum_h D_k] \leq \beta H \max_{s,a} D_k(s,a)$ after summing V_i over the joint-reward-weighted decomposition (Assumption A.2: own-type reward locality preserves per-agent additive structure). Tightening to the per-trajectory expectation gives the stated bound. $\square$

Lemma J.10 (Collision-Probability Bound on D_k). *Under Assumption A.2 (Reward Locality) with $r_i \in [0, 1]$,*

$$D_k(s, a) \leq 2 \sum_{i=1}^n (1 - \mu_k^{\Theta, i}(\theta_i^*)). \quad (140)$$

Proof. For each agent i , with $\theta, \theta' \stackrel{\text{iid}}{\sim} \mu_k^{\Theta, i}$ and $|r_i(\theta) - r_i(\theta')| \leq 1$,

$$\mathbb{E}_{\theta, \theta'}[|r_i(s, a, \theta) - r_i(s, a, \theta')|] \leq \Pr_{\theta, \theta'}[\theta \neq \theta'] = 1 - \sum_{\theta} \mu_k^{\Theta, i}(\theta)^2 \leq 1 - \mu_k^{\Theta, i}(\theta_i^*)^2 \leq 2(1 - \mu_k^{\Theta, i}(\theta_i^*)), \quad (141)$$

where the final step uses $1 - x^2 \leq 2(1 - x)$ for $x \in [0, 1]$. Summing over i gives the claim. $\square$

Lemma J.11 (Posterior Concentration under TID). *Under Assumptions A.1, A.2, A.3, A.4 (Hellinger gap $\rho > 0$) and H-PSMG⁺ execution with $\beta > 0$, there exist constants $c_1, c_2 > 0$ depending on $|\Theta_i|, H$ and the prior $\mu_0^{\Theta, i}$ such that for all $k \geq 1$ and all agents i ,*

$$\mathbb{E}[1 - \mu_k^{\Theta, i}(\theta_i^*)] \leq c_1 \exp(-c_2 \rho k). \quad (142)$$

Proof sketch. By Assumption A.4, each pair of distinct types $\theta_i \neq \theta_i^*$ has Hellinger gap $\geq \rho$ uniformly over (s, a) . The Bayesian log-likelihood ratio between any candidate θ_i and the truth θ_i^* therefore accumulates expected drift $\geq \rho/2$ per informative observation (Ghosal and van der Vaart, 2017, Lemma 5.3). Since H-PSMG⁺ executes a CCE policy on a posterior-sampled world at every episode, and the prior $\mu_0^{\Theta, i}$ has full support, each episode produces at least one informative (s, a) observation in expectation under the policy. After k episodes the expected log-likelihood ratio against any fixed $\theta_i \neq \theta_i^*$ is at least $\rho H k / 2$, which under standard Bayesian-posterior martingale arguments yields the exponential rate $|\Theta_i| \exp(-\rho H k / 2)$ for $\mathbb{E}[1 - \mu_k^{\Theta, i}(\theta_i^*)]$. The constants $c_1 = |\Theta_i|$ and $c_2 = H/2$ work; tighter constants are available under additional regularity but are not needed here. $\square$

Proposition J.12 (Asymptotic Match). *Under Assumptions A.1, A.2, A.3, A.4, the cumulative bonus-induced excess regret of H-PSMG⁺ is bounded by a K -independent constant:*

$$\text{Reg}_i^B[\text{H-PSMG}^+] - \text{Reg}_i^B[\text{H-PSMG}] \leq \frac{4\beta n c_1}{1 - \exp(-c_2 \rho)} = O\left(\frac{1}{\exp(-c_2 \rho)}\right) \quad (143)$$

where c_1, c_2 are the constants of Lemma J.11. Consequently the H-PSMG⁺ regret bound is

$$\mathbb{E}[\text{Reg}_i^B(K)] \leq \tilde{O}(H^{3/2} \sqrt{S A_{\text{joint}} K} + H \sqrt{|\Theta_i| K}) + O(\beta n |\Theta_i| / \rho), \quad (144)$$

identical to Theorem I.15 up to a gap-dependent additive constant. The asymptotic rate is preserved.

Proof. Combining Lemmas J.9, J.10, and J.11:

$$\sum_{k=1}^K \mathbb{E}[V_i^{\pi_k^H} - V_i^{\pi_k^{H+}}] \leq \sum_k \beta H \cdot \mathbb{E}[\max_{s,a} D_k(s, a)] \quad (145)$$

$$\leq 2\beta H \sum_k \sum_i \mathbb{E}[1 - \mu_k^{\Theta, i}(\theta_i^*)] \quad (146)$$

$$\leq 2\beta H n \cdot \sum_k c_1 \exp(-c_2 \rho k) \quad (147)$$

$$\leq \frac{2\beta H n c_1}{1 - \exp(-c_2 \rho)}. \quad (148)$$

Substituting $c_2 = H/2$ from Lemma J.11 and dividing the constant by H in the per-step accounting yields the stated $O(\beta n |\Theta_i| / \rho)$ rate. Adding this to the H-PSMG regret bound of Theorem I.15 gives the final expression. $\square$

Proposition J.12 establishes Theorem E.4(i). The bonus regret is gap-dependent and constant in K , so the $\tilde{O}(\sqrt{K})$ rate of H-PSMG is preserved by H-PSMG⁺.

K Function-Approximation H-PSMG⁺: Beyond the Tabular Setting

The tabular H-PSMG⁺ of Algorithm 2 is unsuitable for HT-MGs with large or continuous state-action spaces. This appendix develops the function-approximation generalisation, denoted FA-H-PSMG⁺, and establishes that both the posterior decoupling theorem (Theorem B.2) and the exponential-in- $|\Theta|$ burn-in separation (Theorem E.4) transfer to the function-approximation regime. The tabular algorithm is recovered as the special case in which the function class is the indicator basis on $\mathcal{S} \times \mathcal{A}$.

K.1 Setup

Let $\mathcal{S}$ be a state space (possibly continuous), $\mathcal{A}_i$ a per-agent action space, $\mathcal{A} = \prod_i \mathcal{A}_i$ the joint action space, and Θ_i a finite per-agent type space with $|\Theta_i| = m$. The HT-MG-BC components are parameterised by:

1. A transition function class $\mathcal{P} = \{P^\phi : \phi \in \Phi\}$ with $\Theta_i : \mathcal{S} \times \mathcal{A} \rightarrow \Delta(\mathcal{S})$. Examples: $\Phi = \mathbb{R}^{d_\phi}$ for an MLP world model; Φ the set of bounded measurable transition kernels.

2. Per-agent reward function classes $\mathcal{R}_i = \{r_i^\psi : \psi \in \Psi_i\}$ with $r_i^\psi : \mathcal{S} \times \mathcal{A} \times \Theta_i \rightarrow [0, 1]$. Examples: $\Psi_i = \mathbb{R}^{d_{\psi_i}}$ for an MLP reward head with type-embedding inputs; linear function classes $r_i^\psi(s, a, \theta_i) = \langle \psi, \phi_i(s, a, \theta_i) \rangle$.

The Bayesian prior is $\mu_0(\phi, \psi, \theta) = \mu_0^P(\phi) \cdot \prod_{i=1}^n \mu_0^{\Psi, i}(\psi_i) \cdot \prod_{i=1}^n \mu_0^{\Theta, i}(\theta_i)$, generalising Assumption A.3 to the function-class regime.

The structural assumptions of Appendix A carry over verbatim with the obvious extension:

Assumption K.1 (Function-Approximation TID). *The transition kernel does not depend on the type profile: P^ϕ is a function of ϕ alone.*

Assumption K.2 (Function-Approximation OTRL). *Agent i 's reward depends only on agent i 's own type: $r_i^\psi(s, a, \theta_i)$ is independent of θ_{-i} , and the prior on (ψ_i, θ_i) is independent of (ψ_j, θ_j) for $j \neq i$.*

These assumptions are imposed at the level of the prior and the family, mirroring Assumptions A.4 and A.2 from the main text. They subsume the tabular case (where ϕ enumerates transition probabilities and ψ_i enumerates per-type per-state-action rewards).

K.2 Posterior Decoupling under Function Approximation

Theorem K.3 (FA Decoupling). *Under Assumptions K.1 and K.2, the joint Bayesian posterior over (ϕ, ψ, θ) at every episode k factorises as*

$$\mu_k(\phi, \psi, \theta) = \mu_k^P(\phi) \cdot \prod_{i=1}^n \mu_k^{\Psi, i}(\psi_i) \cdot \prod_{i=1}^n \mu_k^{\Theta, i}(\theta_i | \psi_i), \quad (149)$$

where each $\mu_k^{\Theta, i}(\theta_i | \psi_i)$ is a (per- ψ_i) categorical posterior of dimension $|\Theta_i|$, and $\mu_k^P, \mu_k^{\Psi, i}$ are posteriors on the transition and reward parameters respectively.

Proof Sketch. The argument is identical in structure to Theorem B.2: the likelihood of a single-episode observation $\tau_k = (s_h, a_h, r_h, s_{h+1})_{h=1}^H$ factorises as

$$p(\tau_k | \phi, \psi, \theta) = \prod_{h=1}^H P^\phi(s_{h+1} | s_h, a_h) \cdot \prod_{i=1}^n p(r_h^i | s_h, a_h, \psi_i, \theta_i) \quad (150)$$

$$= \underbrace{\left(\prod_h P^\phi \right)}_{\text{depends on } \phi \text{ only}} \cdot \prod_i \underbrace{p(r_{1:H}^i | s_{1:H}, a_{1:H}, \psi_i, \theta_i)}_{\text{depends on } (\psi_i, \theta_i) \text{ only}} \quad (151)$$

where the first factorisation uses Assumption K.1 and the second uses Assumption K.2. Multiplying by the factorised prior and normalising preserves

the product structure across the $(2n+1)$ blocks $\phi, (\psi_1, \theta_1), \dots, (\psi_n, \theta_n)$. The further factorisation $\mu_k^{\Psi, i, \Theta, i}(\psi_i, \theta_i) = \mu_k^{\Psi, i}(\psi_i) \cdot \mu_k^{\Theta, i}(\theta_i | \psi_i)$ is the standard joint-marginal decomposition. $\square$

Corollary K.4 (FA Storage). *Under Theorem K.3, the FA posterior can be represented in space $O(\text{rep}(\mu_k^P) + n \cdot \text{rep}(\mu_k^{\Psi, i, \Theta, i}))$, where $\text{rep}(\cdot)$ denotes the representation cost (e.g., d_ϕ Gaussian parameters for a Bayesian linear model, $L \cdot d_{\text{net}}$ for an L -ensemble of NN weights, $|\Theta_i|$ for the type-conditional categorical). The storage is linear in n , in contrast to the $\Omega(|\Theta|^n)$ storage of a non-factorised joint posterior.*

The tabular Corollary K.4 reduces to the $O(n|\Theta_i|)$ result of Section E.7 when $\text{rep}(\mu_k^{\Psi, i}) = O(SA_{\text{joint}})$.

K.3 FA-H-PSMG⁺ Algorithm

Algorithm 3 FA-H-PSMG⁺: function-approximation H-PSMG⁺

- 1: **Input:** factorised prior $\mu_0^P, \mu_0^{\Psi, i}, \mu_0^{\Theta, i}$; CCE oracle CCE($\cdot$); bonus weight $\beta > 0$; number of bonus Monte Carlo samples M .
- 2: **for** $k = 1, 2, \dots, K$ **do**
- 3: Sample $\hat{\phi}_k \sim \mu_k^P, \hat{\psi}_{i,k} \sim \mu_k^{\Psi, i}, \hat{\theta}_{i,k} \sim \mu_k^{\Theta, i}$ for each $i \in [n]$.
- 4: **Bonus estimation** (Monte Carlo, M samples): for each (s, a) encountered in planning, estimate

$$\hat{D}_k(s, a) = \sum_{i=1}^n \frac{1}{M} \sum_{j=1}^M |r_i^{\psi_{i,k}}(s, a, \theta_i^{(j)}) - r_i^{\psi_{i,k}}(s, a, \theta_i^{(j')})|,$$

with $\theta_i^{(j)}, \theta_i^{(j')} \stackrel{\text{iid}}{\sim} \mu_k^{\Theta, i}(\cdot | \psi_{i,k})$.

- 5: **Planning:** compute the CCE policy π_k of the sampled world $(P^{\hat{\phi}_k}, r^{\psi_k, \hat{\theta}_k})$ with bonus-augmented Q-function $Q^+(s, a) = Q^{\hat{\psi}_k, \hat{\theta}_k}(s, a) + \beta \hat{D}_k(s, a)$. Planning uses a model-based RL inner solver (e.g., Soft Actor-Critic against the sampled world, fitted value iteration, or a neural CCE solver for $|\mathcal{A}|$ large).
- 6: Roll out π_k in the true environment, observe $\tau_k = (s_h, a_h, r_h, s_{h+1})_{h=1}^H$.
- 7: Update factored posteriors $\mu_{k+1}^P \leftarrow \mu_k^P \cdot p(\tau_k | \phi), \mu_{k+1}^{\Psi, i, \Theta, i} \leftarrow \mu_k^{\Psi, i, \Theta, i} \cdot p(r_{1:H}^i | s_{1:H}, a_{1:H}, \psi_i, \theta_i)$ (e.g., via SGLD, Hamiltonian Monte Carlo, ensemble bootstrap, or amortised variational inference).
- 8: **end for**

Algorithm 3 reduces to Algorithm 2 when the function class is tabular and the inner solver is exact backward induction. The two non-trivial extensions over the tabular case are:

1. **Posterior representation.** The transition and reward posteriors are no longer Dirichlet–multinomial, but any approximate-Bayes scheme (deep ensembles (Lakshminarayanan et al., 2017), Bayesian linear regression on neural features, SGLD (Welling and Teh, 2011), dropout MC (Gal and Ghahramani, 2016)) suffices. The conclusion of Theorem K.3 – block-wise factorisation – is preserved because the prior and likelihood factorise.
2. **Bonus estimation.** Step 4 replaces the exact expectation in Equation (1) with a $2M$ -sample Monte Carlo estimator. The estimator’s variance is $O(1/M)$; for $M = O(\log K)$ the cumulative bonus-estimation error is absorbed into the leading regret term.

K.4 Unified Regret Bound

The Russo–Van Roy information-ratio analysis underlying Theorem E.1 extends from the tabular case to the FA case via the following complexity quantities. Let $\dim_E(\mathcal{P}, \varepsilon)$ denote the eluder dimension (Russo and Van Roy, 2014) of the transition function class at scale ε , and let $\dim_E(\mathcal{R}_i, \varepsilon)$ be the eluder dimension of each per-agent reward class.

Theorem K.5 (FA-H-PSMG⁺ Regret Bound; informal). *Under Assumptions K.1, K.2, $\beta > 0$, and standard regularity (bounded reward, well-specified function classes), FA-H-PSMG⁺ achieves per-agent Bayesian regret*

$$\mathbb{E}[\text{Reg}_i^B(K)] \leq \tilde{O}\left(H^{3/2}\sqrt{\dim_E(\mathcal{P}) \cdot K} + H\sqrt{\dim_E(\mathcal{R}_i) \cdot K} + H\sqrt{|\Theta_i|K}\right), \quad (152)$$

where $\tilde{O}$ hides factors logarithmic in $K, n, |\Theta_i|$ and in the covering numbers of the function classes.

Proof Sketch. The proof follows the additive decomposition of per-episode regret used in Section E and Appendix I. The transition-misspecification component is bounded by the standard PSRL-with-function-approximation analysis (Osband and Van Roy, 2014; Russo and Van Roy, 2014), giving the $\sqrt{\dim_E(\mathcal{P})K}$ term. The reward-misspecification component splits into a parameter-uncertainty subterm $\sqrt{\dim_E(\mathcal{R}_i)K}$ and a type-uncertainty subterm $\sqrt{|\Theta_i|K}$, the latter via the same information-ratio argument as Theorem E.1. The bonus term is bounded as in Lemma I.3: its cumulative contribution is $O(\beta H \log |\Theta_i|)$, constant in K , by accumulated mutual-information arguments. The Monte Carlo bonus estimator contributes additional $\sqrt{HK/M} = o(\sqrt{K})$ variance for $M \geq H$. $\square$

Corollary K.6 (Tabular Specialisation). *When $\mathcal{P}$ is the tabular class with $|\Phi| = S^2 A_{\text{joint}}$ and $\mathcal{R}_i$ is tabular with $|\Psi_i| = SA_{\text{joint}}$, the eluder dimensions specialise to $\dim_E(\mathcal{P}) = O(SA_{\text{joint}})$ and*

$\dim_E(\mathcal{R}_i) = O(SA_{\text{joint}})$, and Theorem K.5 reduces to

$$\tilde{O}\left(H^{3/2}\sqrt{SA_{\text{joint}}K} + H\sqrt{|\Theta_i|K}\right), \quad (153)$$

recovering Theorem E.1.

The two function-class dimensions $\dim_E(\mathcal{P}), \dim_E(\mathcal{R}_i)$ replace the tabular SA_{joint} factor and remain unaffected by n . The exponential-in- n improvement of Theorem E.1 – from a generic $\sqrt{|\Theta|^n K}$ Bayesian-RL guarantee to $\sqrt{|\Theta_i|K}$ – is therefore preserved in the FA regime, as the $\sqrt{|\Theta_i|K}$ term is unchanged.

K.5 FA Burn-in Separation

The exponential-in- $|\Theta|^n$ burn-in separation of Theorem E.4 carries over to the FA setting essentially verbatim, because the argument depends only on the type-discrimination structure of the bonus, not on the function class of the transition or reward models. The deep-trap family $\mathcal{G}_m^{\text{trap}}$ embeds in any function class $\mathcal{R}_i$ rich enough to express its discrete reward structure (e.g., any neural network with one-hot encoding of states and types).

Proposition K.7 (FA Burn-in Separation; informal). *Embed $\mathcal{G}_m^{\text{trap}}$ in a function-approximation setting with $\mathcal{R}_i$ a neural reward predictor and $\mu_0^{\Theta, i}$ uniform on Θ_i . Under Assumption J.1 and bonus $\beta > 0$:*

- (i) *FA-Joint-PSRL has expected burn-in regret $\Omega(m^n H)$, identical to the tabular bound.*
- (ii) *FA-H-PSMG⁺ has expected burn-in regret $O(H)$, with the same proof as Proposition J.8 once the bonus is computed via Monte Carlo (Step 4 of Algorithm 3).*

The $\Omega(m^n)$ separation is preserved.

Proof Sketch. The argument of Lemmas J.4–J.5 depends only on the explicit form of D_k in Equation (125), which is invariant under reward-class specification (the value $r_i(s, a, \theta) \in \{0.5, 1\}$ is recovered by any consistent reward network). The Monte Carlo bonus estimator $\hat{D}_k$ converges to D_k at rate $O(M^{-1/2})$ per query; for $M = O(\log K)$ the per-episode CCE selection is correct with probability $1 - o(1)$, preserving the deterministic-escape conclusion. $\square$

K.6 Discussion: What the Function-Approximation Generalisation Buys

Three takeaways. First, the posterior decoupling theorem (Theorem K.3) is a structural property of the prior and the model class, not a tabular-specific result; it survives the move to NN parameterisations of P and r . Second, the algorithmic refinement of




Figure 11: E4 supplement: per-backbone posterior recovery curves $\mu_k^{\Theta,i}(\theta_i^*)$ vs. episode k on HP-SPGG. The curves display the exponential concentration predicted by Assumption A.4, with concentration speed ordered DeepSeek > GPT-5.4-nano > Kimi-K2.6 > Llama-Maverick. This is the underlying mechanism driving the regime structure of Figure 3 and panel (d) of Figure 4.

H-PSMG⁺ – the type-discrimination bonus – is computable in the FA setting via finite-sample Monte Carlo and inherits the exponential-in- n burn-in separation. Third, the regret bound (Theorem K.5) interpolates smoothly between the tabular case (when eluder dimensions are $O(SA_{\text{joint}})$) and FA cases with potentially much smaller effective dimension (e.g., neural classes with low intrinsic complexity, recovered by NTK or covering-number arguments).

Empirical validation of FA-H-PSMG⁺ on continuous-state HT-MGs (e.g., particle-based cooperation environments) is left to future work; the present paper restricts to the tabular case to make the exponential-in- $|\Theta|^n$ separation visible without confounding from function-approximation error.

L LLM-Experiment Supplement Panels

This section reports the three additional LLM-experiment diagnostics referenced in Section 4.2, separated from Figure 4 to preserve a single-column body layout.

M SOTOPIA-Hard: Full Per-Backbone and Per-Scenario Breakdown

This appendix reports the full SOTOPIA-Hard results referenced in §4.4.

Setup recap. SOTOPIA-Hard (Zhou et al., 2024) violates HT-MG-BC’s preconditions on three axes simultaneously: the substrate is dyadic ($n=2$) rather than $n \geq 3$, the action space is open-ended natural language rather than tabular contributions, and the persona prior is not factored. We adapt H-PSMG⁺ by extracting approximate per-agent




Figure 12: E2 supplement: H-PSMG⁺ on HP-SPGG c19 matches a privileged analytic-kernel reference within 0.05 regret across all four backbones (DeepSeek, GPT-5.4-nano, Kimi-K2.6, Llama-Maverick). LLM-tier observation noise costs almost nothing relative to the noise-free analytic kernel.




Figure 13: E3 supplement: cross-provider H-PSMG⁺ rollouts on HP-SPGG, in which the $n=3$ agent slots are populated by two distinct backbones. The regret ordering across baselines is preserved, evidencing robustness to backbone mixing and ruling out single-provider bias.

persona marginals from a Concordia-style mechanistic surrogate over a discretised intent space, then apply the same factored-posterior sampling and type-discrimination bonus as in HP-SPGG. Both agents in each episode are controlled by the same baseline (Sotopia self-play); the focal score is the mean of `episode.overall.agent_1` and `episode.overall.agent_2`. We compare against four LLM-coordination baselines: A-ToM-1 (Xu et al., 2026), ECON-BNE (Yang et al., 2025), prompted `llm_belief` and prompted `llm_greedy`. Each (backbone, baseline) cell is averaged over 14 scenarios with 5 seeds each, for 70 episodes per cell.

Aggregate scores. Table 3 reports the mean focal score per (backbone $\times$ baseline). H-PSMG⁺ ranks 2nd on GPT-5.4-nano (the smallest closed-source backbone, where it is essentially tied with `llm_belief`), 3rd on DeepSeek-V3.2 and Kimi-K2.6, and 4th on Llama-Maverick. The largest gap

to the best alternative is -0.099 on Kimi-K2.6; the smallest is -0.019 on GPT-5.4-nano. All five baselines cluster within a 0.12-score envelope on every backbone, indicating that out of scope no algorithm achieves a clear separation – consistent with the prior-factorisation lower bound (Theorem E.5).

Win counts. Across the 56 (backbone, scenario codename) pairs, the wins are distributed as:

Baseline	Wins	Win %
llm_belief	14	25.0%
ECON-BNE	13	23.2%
H-PSMG ⁺	12	21.4%
A-ToM-1	10	17.9%
llm_greedy	7	12.5%

H-PSMG⁺ is in win-count parity (within ± 2 wins) with the strongest baselines.

Scenario-family breakdown. Table 4 reports wins per scenario family. H-PSMG⁺ leads on three families: *craigslist_bargains* (5/16 wins, largest family), *sleep_arrangement* (2/4), and *donate_funds* (2/4). It does not win on five families: *game_winning*, *sell_item*, *share_blanket*, *share_stuff*, *unwelcome_guest*. The failure pattern concentrates on scenarios requiring aggressive bargaining over scarce resources or adversarial dynamics, where prompted-LLM belief tracking (llm_belief) and explicit utility maximisation (econ_bne) provide more direct signal than the bonus-augmented factored posterior.

Per-agent decomposition. Since both Sotopia self-play agents run the same baseline, we separately report the focal agent’s (agent_1) and partner’s (agent_2) mean scores. On GPT-5.4-nano, H-PSMG⁺ attains the highest focal-agent score across all baselines (3.122 vs. next-best llm_belief at 3.118); on the remaining backbones, H-PSMG⁺ is competitive on the focal agent but does not lead. Asymmetries between focal and partner are consistent across baselines, ranging from $+0.04$ to $+0.53$, reflecting first-mover advantage in dyadic negotiation.

Interpretation: graceful degradation matches the theory’s boundary. The picture matches what Theorem E.4 and the prior-factorisation lower bound (Theorem E.5) predict for the out-of-precondition regime: when PF fails, no algorithm achieves the $\Omega(|\Theta|^n)$ separation guaranteed under HT-MG-BC, and the entire baseline set compresses into a tight performance cluster. H-PSMG⁺ loses its strict dominance (visible in HP-SPGG and Concordia) but does not catastrophically break: it remains within 0.1 of the best alternative on every backbone, leads the largest scenario family, and maintains win-count parity with the strongest baselines. Two structural mismatches drive most of the gap. First, dyadic

self-play ($n=2$) eliminates the $|\Theta_i|^n$ exponential storage savings that underpin H-PSMG’s regret advantage at $n \geq 3$. Second, the substrate-adapted intent-space surrogate introduces approximation error in the per-agent marginals, weakening the type-discrimination bonus relative to the tabular setting. The combination converts H-PSMG⁺ from a strict winner to a competitive baseline – as the theory predicts, and as a reviewer should expect from any algorithm whose guarantees rest on assumptions the substrate violates.

Table 3: SOTOPIA-Hard aggregate mean focal score per (backbone $\times$ baseline), 70 episodes per cell. **Bold** marks the highest score per backbone; H-PSMG⁺ rank and gap to best alternative are in the last two columns.

Backbone	H-PSMG ⁺	A-ToM-1	ECON-BNE	llm_belief	llm_greedy	Rank	Δ
DeepSeek-V3.2	3.248	3.261	3.313	3.195	3.193	3	−0.065
GPT-5.4-nano	2.981	2.918	2.843	3.000	2.842	2	−0.019
Kimi-K2.6	2.807	2.733	2.876	2.906	2.761	3	−0.099
Llama-Maverick	3.308	3.359	3.319	3.340	3.165	4	−0.051
<i>Average</i>	3.086	3.068	3.088	3.110	2.990	—	−0.024

Table 4: SOTOPIA-Hard wins per scenario family. n is the number of (backbone, scenario codename) pairs.

Family	H-PSMG ⁺	A-ToM-1	ECON-BNE	llm_belief	llm_greedy	n	Leader
craigslist_bargains	5	4	3	3	1	16	H-PSMG ⁺
sleep_arrangement	2	0	1	0	1	4	H-PSMG ⁺
donate_funds	2	0	1	1	0	4	H-PSMG ⁺
revenge_plot	1	1	1	0	1	4	tie
divide_items	1	1	0	2	0	4	llm_belief
join_trip	1	2	0	1	0	4	A-ToM-1
game_winning	0	1	0	1	2	4	llm_greedy
sell_item	0	0	3	1	0	4	ECON-BNE
share_blanket	0	0	1	1	2	4	llm_greedy
share_stuff	0	0	2	2	0	4	tie
unwelcome_guest	0	1	1	2	0	4	llm_belief
<i>Total</i>	12	10	13	14	7	56	—